x86 Initial Booting Sequence

Advanced Operating Systems and Virtualization
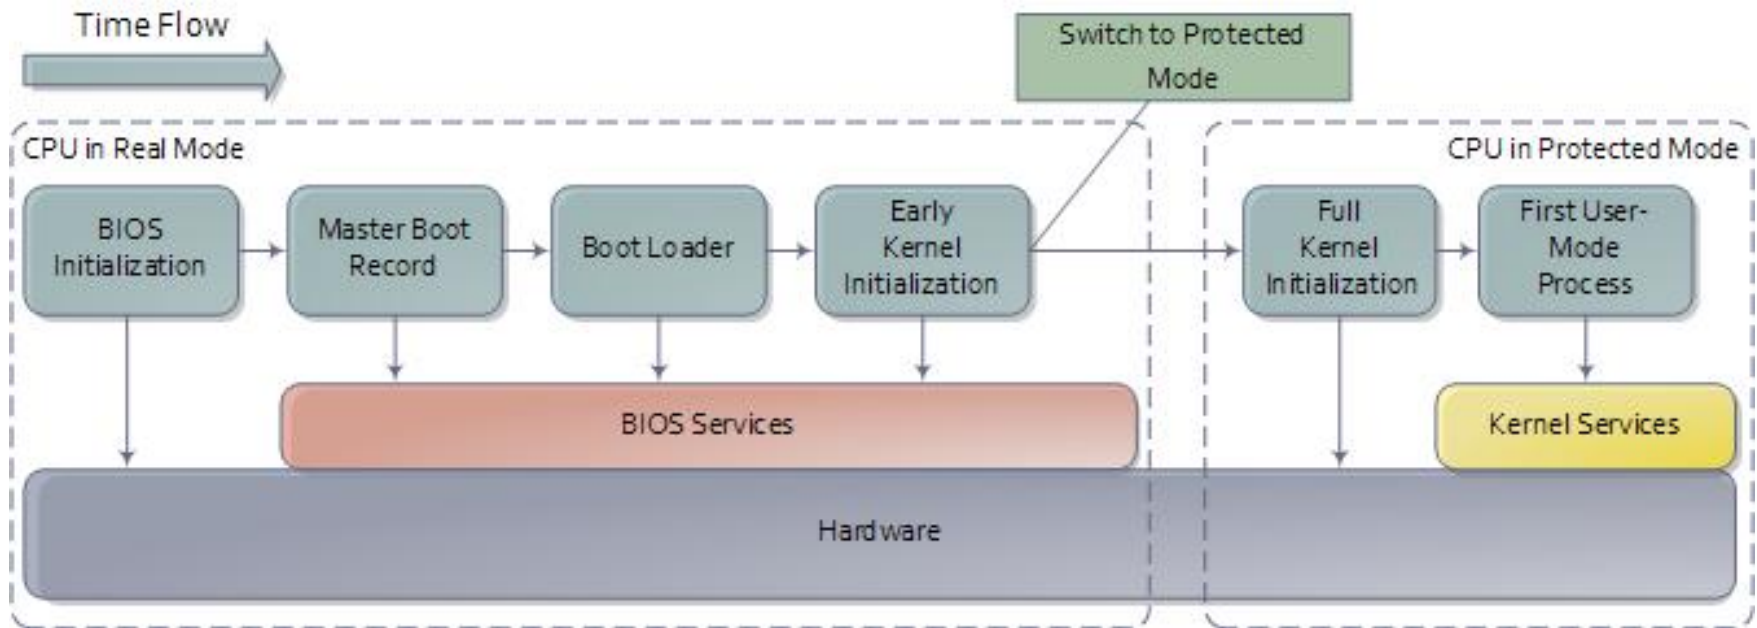Alessandro Pellegrini
A.Y. 2017/2018

The whole sequence at a glance

Boot Sequence

BIOS/UEFI	The actual Hardware Startup
Bootloader Stage 1	Executes the Stage 2 bootloader (skipped in case of UEFI)
Bootloader Stage 2	Loads and starts the Kernel
Kernel Startup	The Kernel takes control of and initializes the machine (machine−dependent operations)
Init	First process: basic environment initialization (e.g., SystemV Init, systemd)
Runlevels/Targets	Initializes the user environment (e.g., single−user mode, multiuser, graphical, …)

Hardware Power Sequences: The Pre-Pre-Boot

- When someone pushes the power button, the CPU can't simply jump up and start fetching code from flash memory

- The hardware waits for the power supply to settle to its nominal state

- Additional voltages must be supplied:
  - On x86 systems: 1.5, 3.3, 5, and 12 V
  - Power Sequencing: these must be provided in a particular order

Hardware Power Sequences: The Pre-Pre-Boot

- The power is sequenced by controlling analog switches, typically field-effect transistors

- The sequence is often driven by a Complex Program Logic Device (CPLD)

- Platform clocks are derived from a small number of input clock and oscillator sources.
  - The devices use phase-locked loop circuitry to generate the derived clocks used for the platform.
  - These clocks take time to converge.

Hardware Power Sequences: The Pre-Pre-Boot
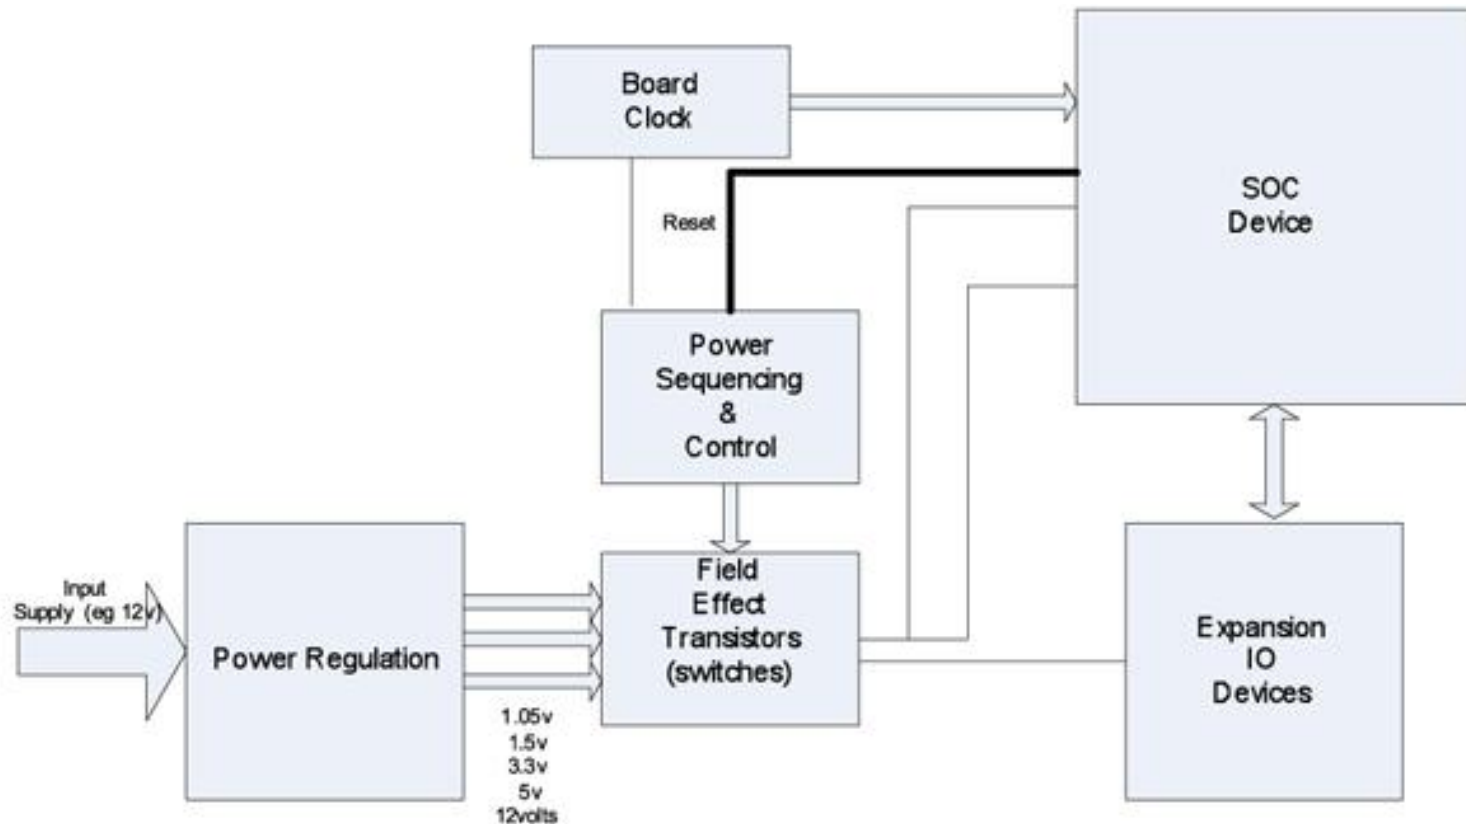
Initial life of the System

- The power-sequencing CPLD can de-assert the reset line to the processor

- At this point, the system is in a very *basic state*:
  - Caches are disabled
  - The Memory Management Unit (MMU) is disabled
  - The CPU is executing in **Real Mode** (8086-compatible)
  - Only one core can run actual code
  - Nothing is in RAM (what to execute?)

Segmented Memory

Address space

Segment B
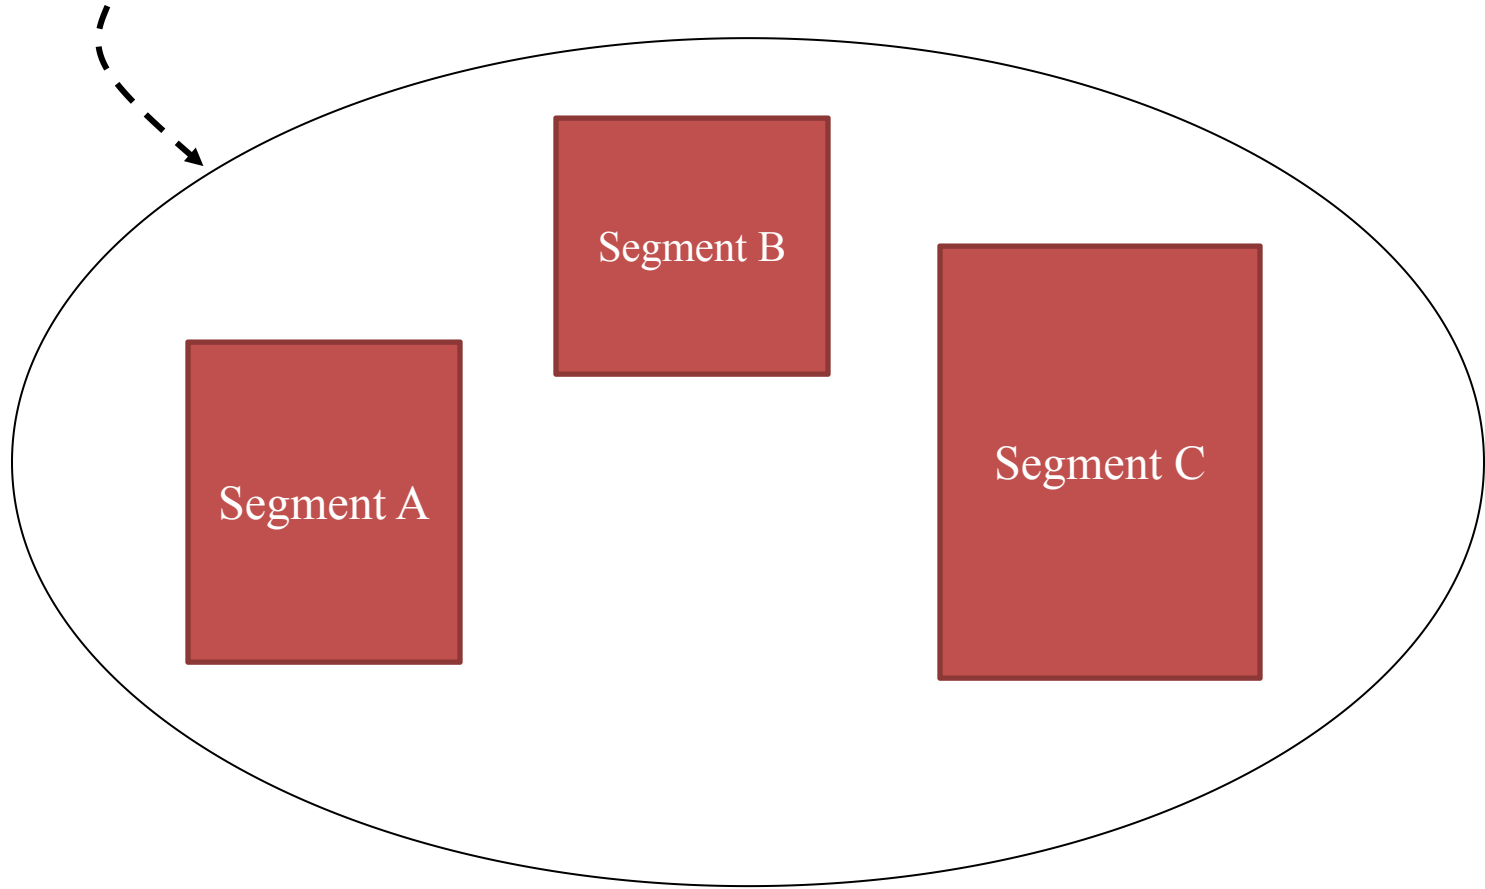
Segment C

Segment A

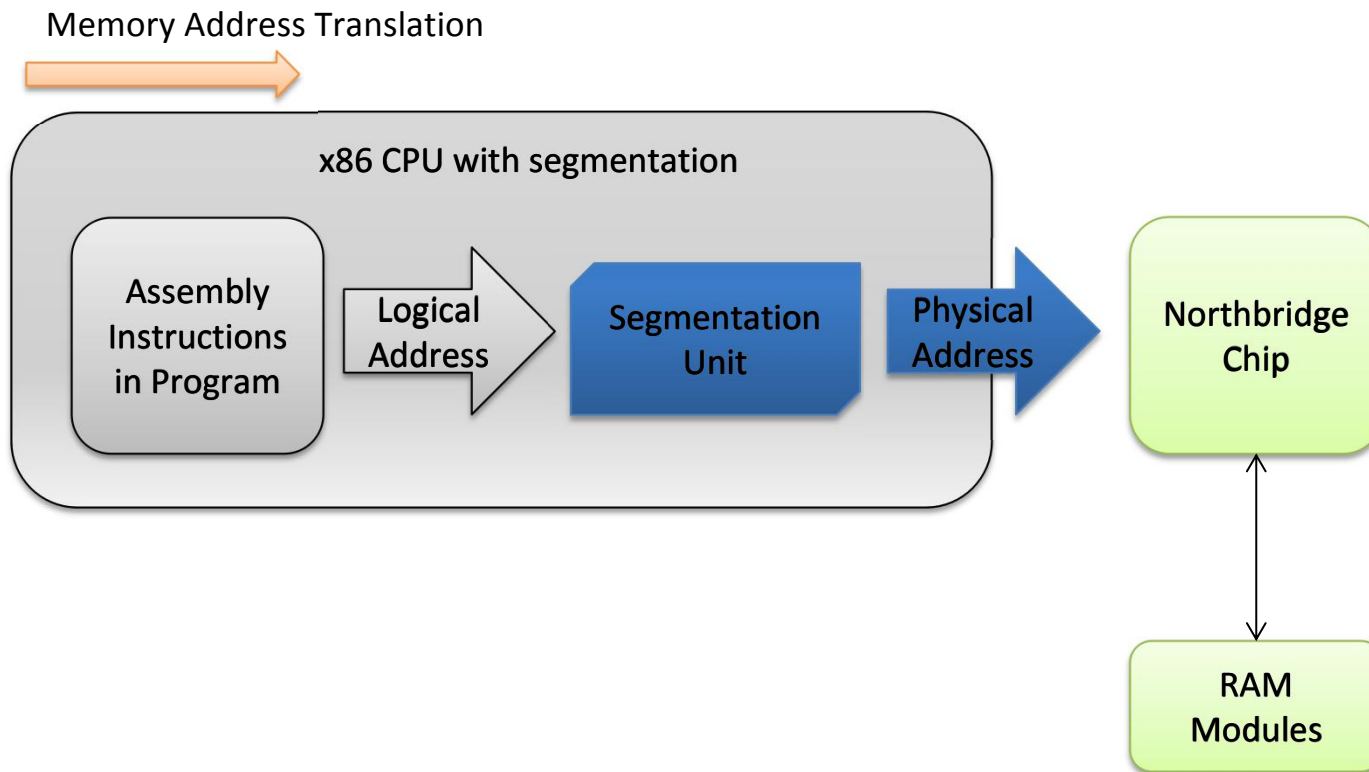*logical address = <seg.id : offset> (es. <A : 0x10>)*

Segmentation-based addressing

- There are **4 basic 16-bit segment registers**:
  - CS: code segment
  - DS: data segment
  - SS: stack segment
  - ES: extra segment (to be used by the programmer)

- Intel 80386 (1985) added two new registers
  - FS and GS, with no predefined usage

Segmentation-based addressing

- The CPU resolves addresses as:

Segmentation Nowadays

- Segmentation is still present and always enabled

- Each instruction that touches memory implicitly uses a segment register:
  - a `jump` instruction uses CS
  - a `push` instruction uses SS

- Most segment registers can be loaded using a `mov` instruction

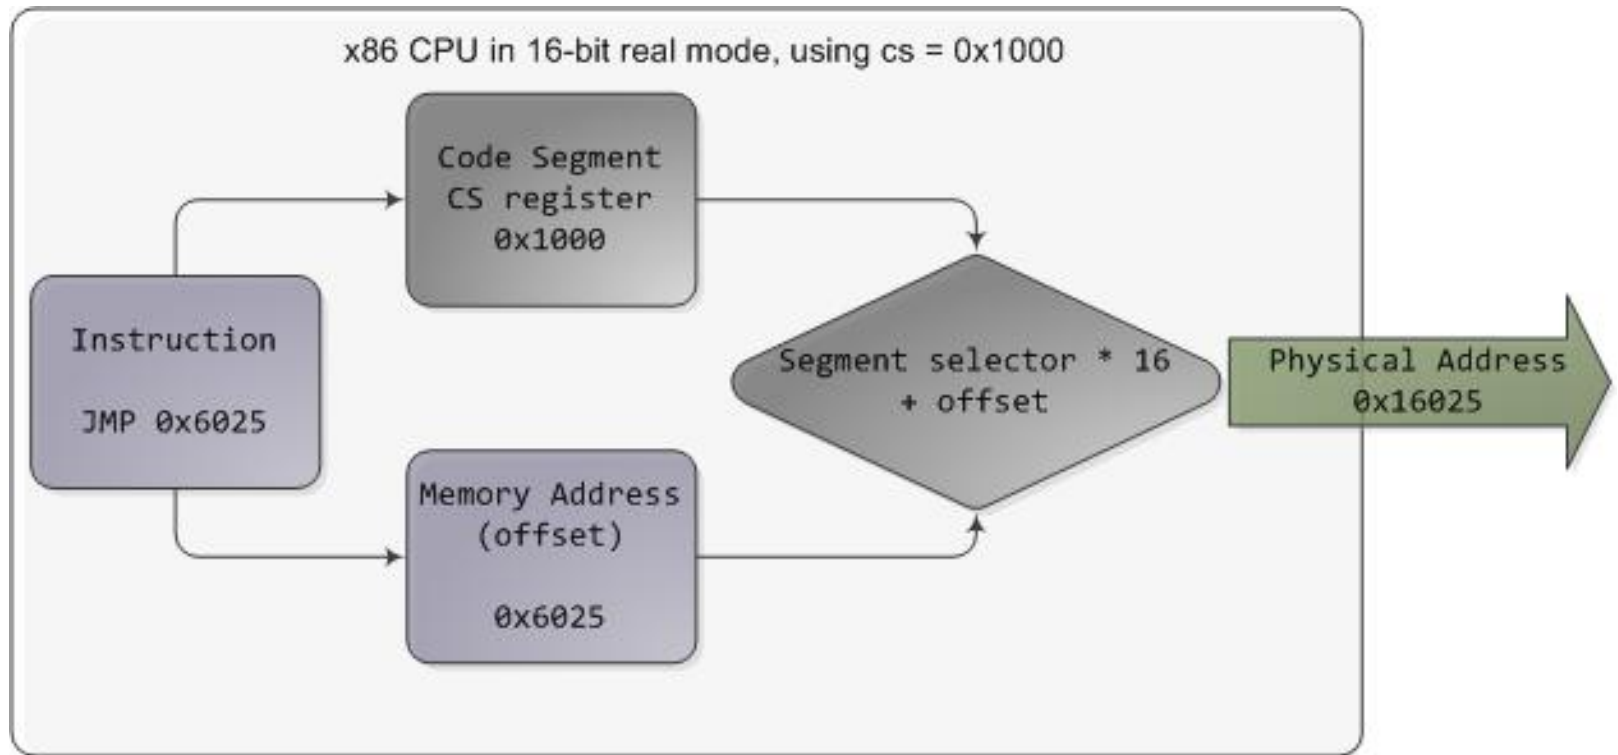- CS can be loaded only with a `jmp` or a `call`
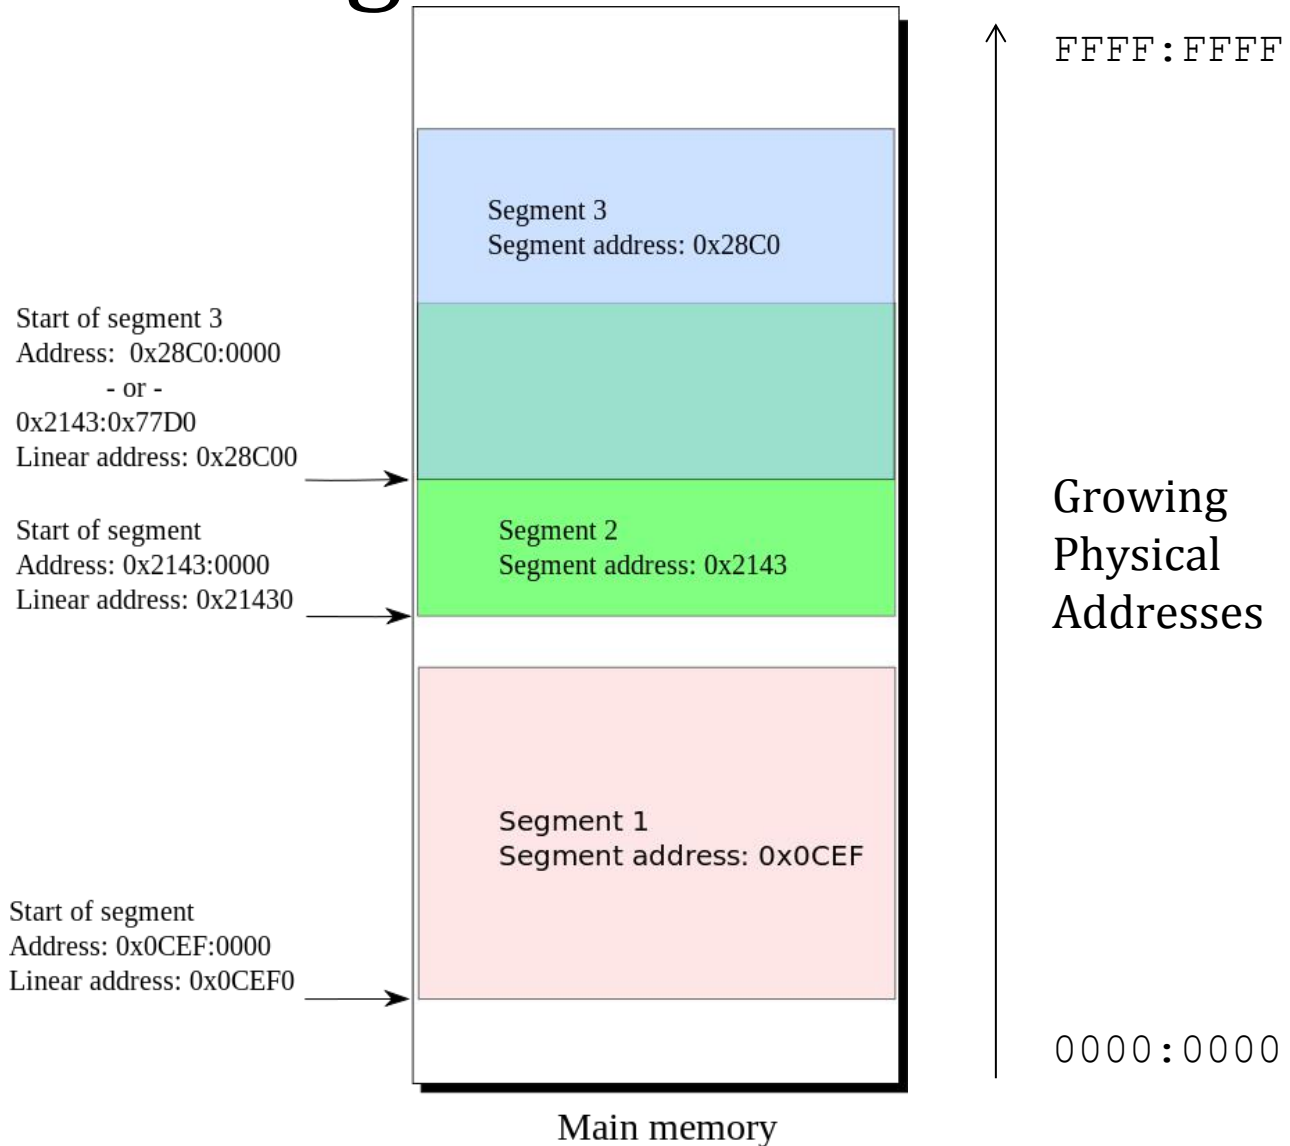
x86 Real Mode

- 16-bit instruction execution mode
- 20-bit segmented memory address space
  - 1 MB of total addressable memory
- Address in segment registers is the 16-bits higher part
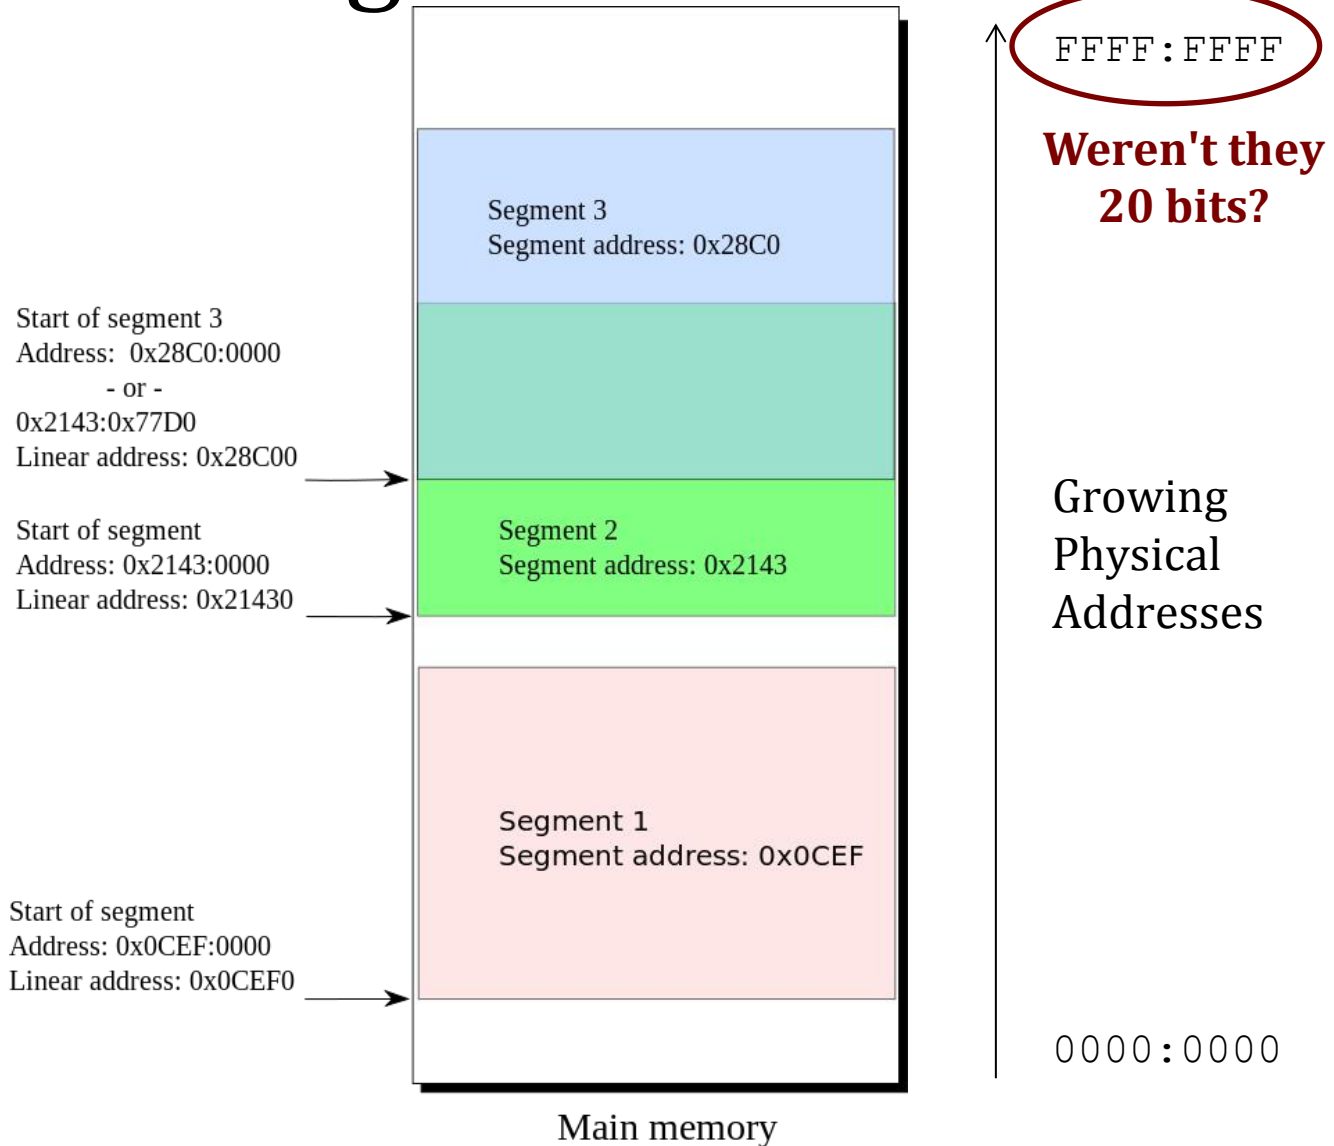- Each segment can range from 1 byte to 65,536 bytes (16-bit offset)
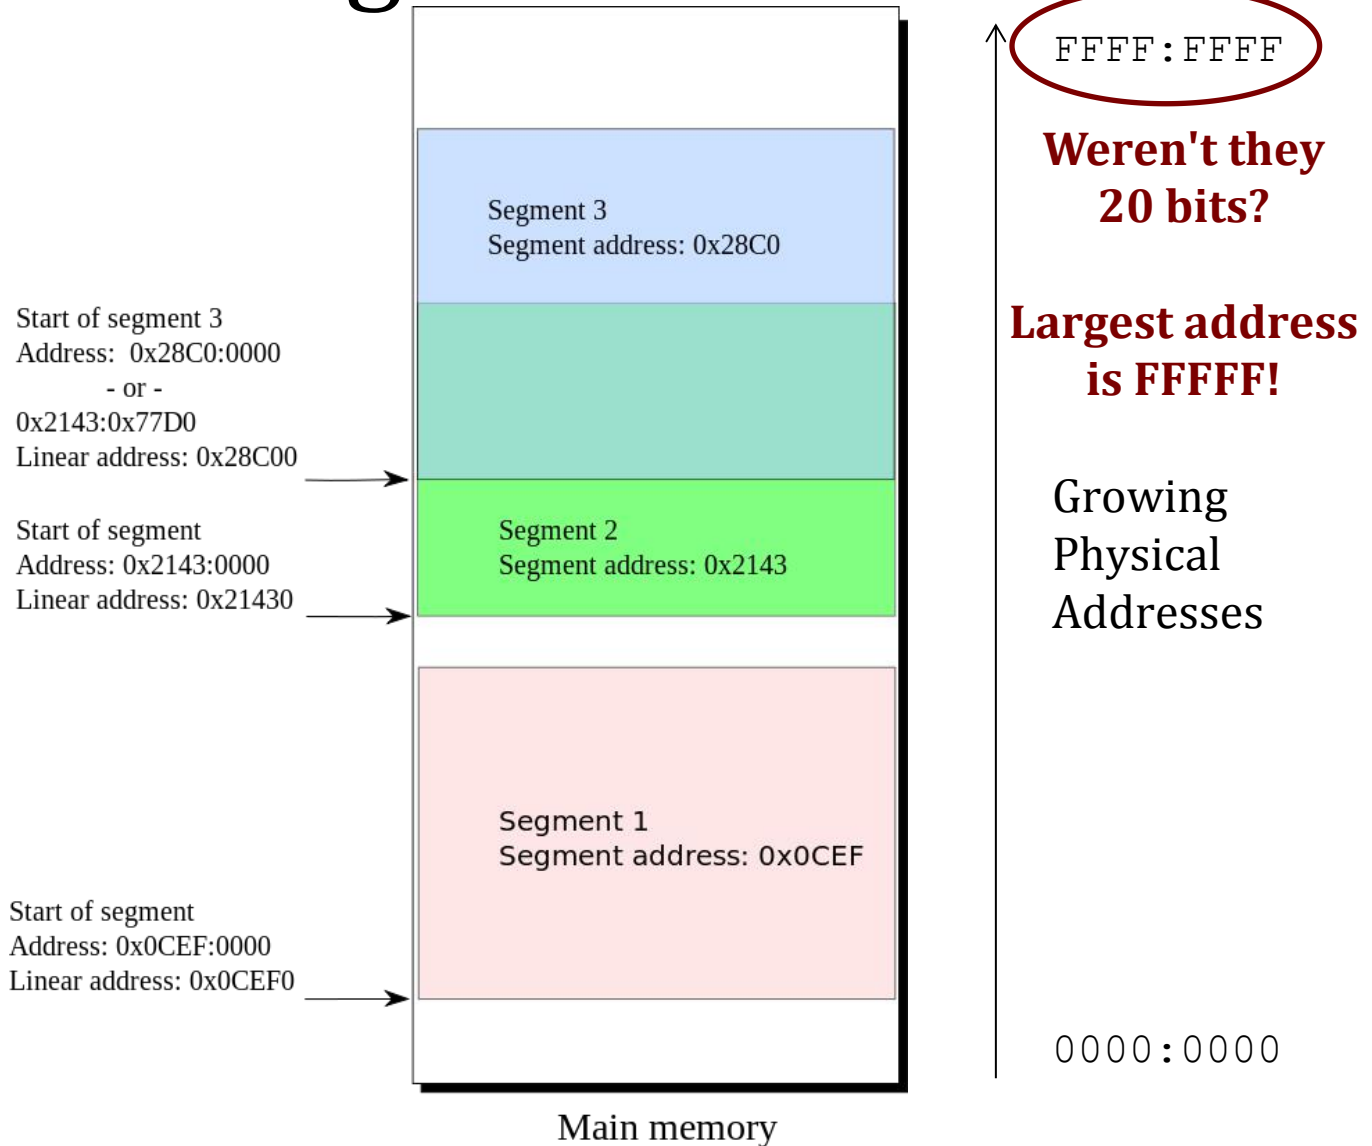
Real Mode Addressing Resolution

Addressing in x86 Real Mode

FFFF:FFFF

Segment 3
Segment address: 0x28C0

Start of segment 3
Address: 0x28C0:0000
- or -
0x2143:0x77D0
Linear address: 0x28C00

Start of segment
Address: 0x2143:0000
Linear address: 0x21430

Segment 2
Segment address: 0x2143

Growing
Physical
Addresses

Segment 1
Segment address: 0x0CEF

Start of segment
Address: 0x0CEF:0000
Linear address: 0x0CEF0

0000:0000

Main memory

Addressing in x86 Real Mode

FFFF:FFFF

Weren't they 20 bits?

Segment 3
Segment address: 0x28C0

Start of segment 3
Address: 0x28C0:0000
- or -
0x2143:0x77D0
Linear address: 0x28C00

Start of segment
Address: 0x2143:0000
Linear address: 0x21430

Segment 2
Segment address: 0x2143

Growing
Physical
Addresses

Segment 1
Segment address: 0x0CEF

Start of segment
Address: 0x0CEF:0000
Linear address: 0x0CEF0

0000:0000

Main memory

Addressing in x86 Real Mode

Segment 3
Segment address: 0x28C0

Start of segment 3
Address: 0x28C0:0000
   - or -
0x2143:0x77D0
Linear address: 0x28C00

Start of segment
Address: 0x2143:0000
Linear address: 0x21430

Segment 2
Segment address: 0x2143

Segment 1
Segment address: 0x0CEF

Start of segment
Address: 0x0CEF:0000
Linear address: 0x0CEF0

Main memory

FFFF:FFFF

Weren't they 20 bits?

Largest address is FFFFF!

Growing
Physical
Addresses

0000:0000

First Fetched Instruction

- The PC runs in a *quasi*-real mode
  - Bit 20 is shortcircuited to 1
- The first fetched address is `F000:FFF0`
  - This is known as the *reset vector*
  - On IBM PCs this is mapped to a ROM: the **BIOS**
  - This gives space only to 16 bytes from the top of ROM memory:
    ```
    ljmp $0xf000,$0xe05b
    ```
- This is where the BIOS code is loaded

BIOS Operations

- The BIOS first looks for video adapters that may need to load their own routines
  - These ROMs are mapped from C000:0000 to C780:0000

- Power-On Self-Test (POST)
  - Depends on the actual BIOS
  - Often involves testing devices (keyboard, mouse)
  - Video Card Initialization
  - RAM consistency check

BIOS Operations

- Boot configuration loaded from CMOS (64 bytes)
  - For example, the *boot order*
- Shadow RAM initialization
  - The BIOS copies itself into RAM for faster access
- The BIOS tries to identify the Stage 1 bootloader, (512 bytes) using the specified boot order and loads it to memory at `0000:7c00`
- Control is given with:

  ```
  ljmp $0x0000,$0x7c00
  ```

The RAM after the BIOS startup

```
                        ←——————————— 0x00100000 (1 Mb)

      BIOS ROM

                        ←——————————— 0x000F0000 (960 Kb)

   16-bit devices,
   expansion ROM

                        ←——————————— 0x000C0000 (768 Kb)

    VGA Display

                        ←——————————— 0x000A0000 (640 Kb)

                            The bootloader
                            is loaded here
                                              The only available
    Low Memory                                "RAM" in the
                                              early days

                        ←——————————— 0x00000000
```

Boot Sequence

BIOS/UEFI	The actual Hardware Startup
Bootloader Stage 1	Executes the Stage 2 bootloader (skipped in case of UEFI)
Bootloader Stage 2	Loads and starts the Kernel
Kernel Startup	The Kernel takes control of and initializes the machine (machine–dependent operations)
Init	First process: basic environment initialization (e.g., SystemV Init, systemd)
Runlevels/Targets	Initializes the user environment (e.g., single–user mode, multiuser, graphical, ...)

The Boot Sector

- The first device sector keeps the so called Master Boot Record (MBR)
- This sector keeps executable code and a 4-entry partition table, identifying different device partitions (in terms of its positioning on the device)
- In case the partition is extended, then it can additionally keep up to 4 sub-partitions (extended partition)

The Device Organization

Boot code

Partition table

Boot partition

Partition 1

Partition 3 (extended)

Boot sector: it can contain additional boot code

Extended partition boot record
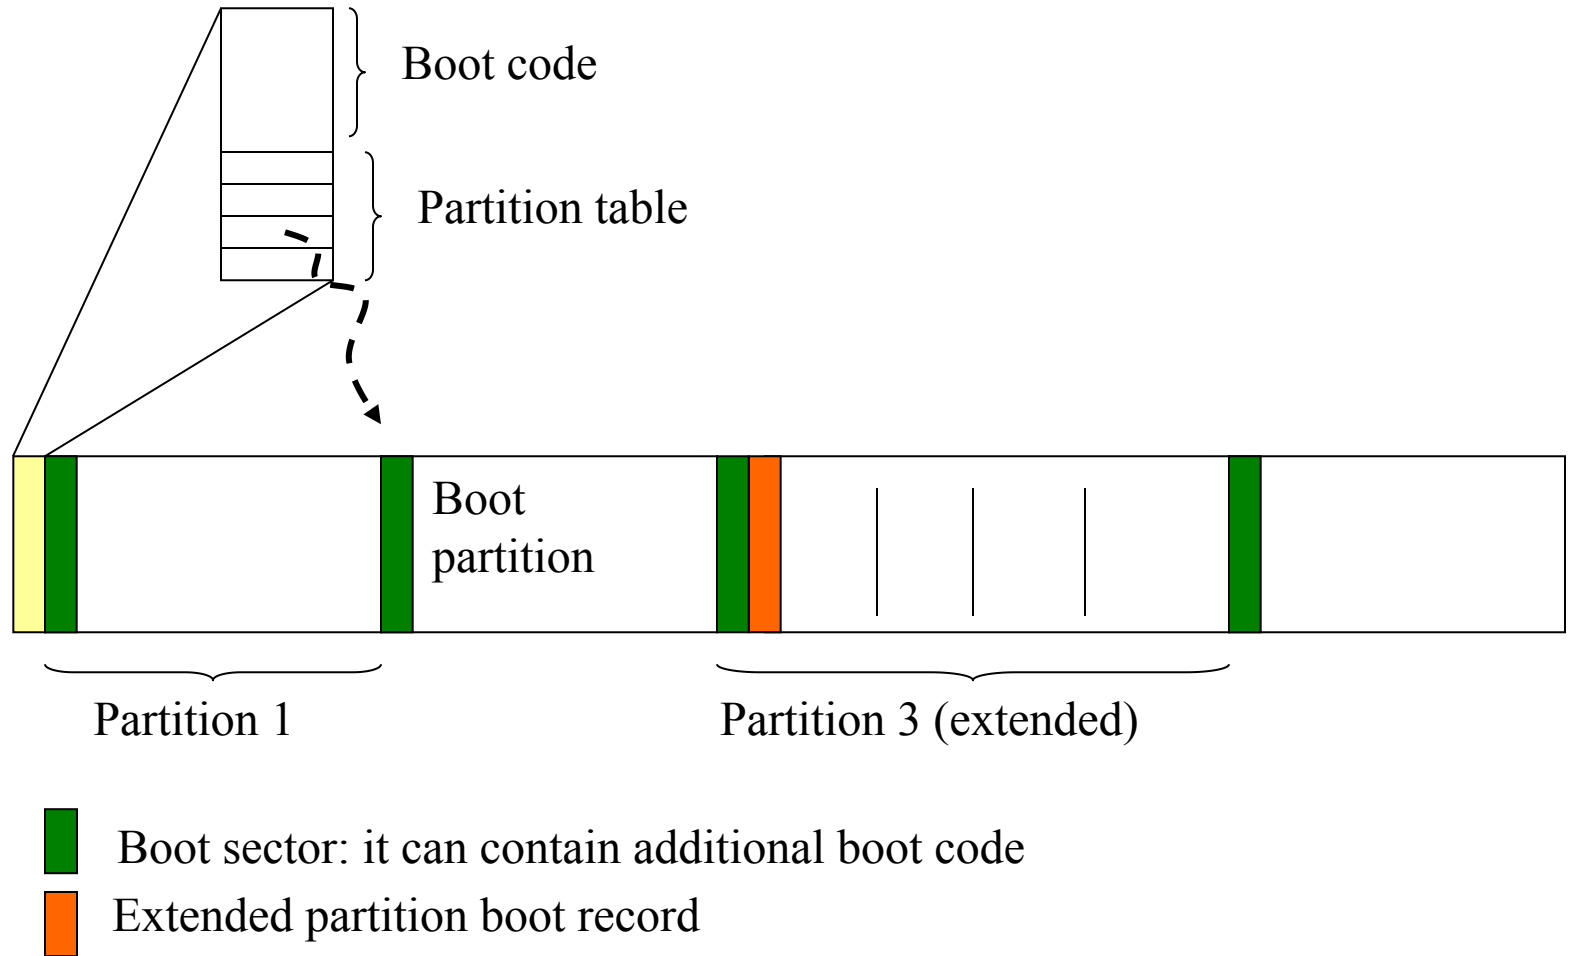
The Master Boot Record (MBR)

- This implements the Stage 1 bootloader
- (Less than) 512 bytes can be used to load the operating system

Offset	Size (bytes)	Description
0	436 (to 446, if you need a little extra)	MBR **Bootstrap** (flat binary executable code)
0x1b4	10	Optional "unique" disk ID[1]
0x1be	64	MBR **Partition Table**, with 4 entries (below)
0x1be	16	First partition table entry
0x1ce	16	Second partition table entry
0x1de	16	Third partition table entry
0x1ee	16	Fourth partition table entry
0x1fe	2	(0x55, 0xAA) "Valid bootsector" signature bytes

The Master Boot Record (MBR)

- The initial bytes of the MBR can contain the *BIOS Parameter Block* (BPB)

- It is a data structure describing the physical layout of a data storage volume

  - It is used, e.g., by FAT16, FAT32, and NTFS

- This eats up additional space, and must be placed *at the beginning* of the MBR!

  - How to execute the code?

The Master Boot Record (MBR)

```
.code16
 .text
.globl _start;

_start:
jmp .stage1_start

OEMLabel:  .string "BOOT"
BytesPerSector:  .short 512
SectorsPerCluster: .byte 1
ReservedForBoot: .short 1
NumberOfFats:  .byte 2
RootDirEntries:  .short 224
LogicalSectors:  .short 2880
MediumByte:  .byte 0x0F0
SectorsPerFat:  .short 9
SectorsPerTrack: .short 18
```

```
Sides:    .short 2
HiddenSectors:  .int 0
LargeSectors:  .int 0
DriveNo:  .short 0
Signature:  .byte 41 #41 = floppy
VolumeID:  .int 0x00000000
VolumeLabel:  .string "myOS"
FileSystem:  .string "FAT12"

 .stage1_start:
    cli # Disable interrupts
    xorw %ax,%ax # Segment zero
    movw %ax,%ds
    movw %ax,%es
    movw %ax,%ss
    ...
```

The Master Boot Record (MBR)

```
.code16
 .text
.globl _start;

_start:
jmp .stage1_start

OEMLabel:  .string "BOOT"
BytesPerSector:  .short 512
SectorsPerCluster: .byte 1
ReservedForBoot: .short 1
NumberOfFats:  .byte 2
RootDirEntries:  .short 224
LogicalSectors:  .short 2880
MediumByte:  .byte 0x0F0
SectorsPerFat:  .short 9
SectorsPerTrack: .short 18
```

```
Sides:    .short 2
HiddenSectors:  .int 0
LargeSectors:  .int 0
DriveNo:  .short 0
Signature:  .byte 41 #41 = floppy
VolumeID:  .int 0x00000000
VolumeLabel:  .string "myOS"
FileSystem:  .string "FAT12"

 .stage1_start:
   cli # Not safe here!
   xorw %ax,%ax # Segment zero
   movw %ax,%ds
   movw %ax,%es
   movw %ax,%ss
   ...
```

What about CS?

The Stage 1 Bootloader must…

- Enable address A20

- Switch to 32-bit **protected mode**

- Setup a stack

- Load the kernel
  - Yet, the kernel is on disk: how to navigate the file system? There is not much space for code…
  - Load the Stage 2 bootloader!

A20 Enable

- Intel 80286 increased the addressable memory to 16 Mb (24 address lines)

- How to keep backward compatibility with 8086?
  - "wrap-around" problem
  - By default address line 20 is forced to zero!

- How to enable/disable this line?
  - Use the 8042 keyboard controller (sic!)
  - It had a spare pin which they decided to route the A20 line through

A20 Enable

- The output port of the keyboard controller has a number of functions.
- Bit 1 is used to control A20:
  - 1 = enabled
  - 0 = disabled
- Port 0x64 is used to "communicate" an operation to the controller
  - 0xd1 means "write"
- 0xdd and 0xdf enable/disable A20, when sent to port 0x60
  - You have to wait for previous operations to complete (the controller is slow)

A20 Enable

```
call wait_for_8042
movb $0xd1, %al #command write
outb %al, $0x64
call wait_for_8042
movb $0xdf, %al # Enable A20
outb %al, $0x60
call wait_for_8042
...

wait_for_8042:
  inb %al, $0x64
  tesb $2, %al     # Bit 2 set = busy
  jnz wait_for_8042
  ret
```

x86 Protected Mode

- This execution mode was introduced in 80286 (1982)
- With 80386 (1985) it was extended by adding **paging**
- CPUs start in Real Mode for backwards compatibility
- Still today, x86 Protected Mode must be activated during system startup

x86_64 Registers

Registri Fisici e Registri Virtuali

AL AX EAX RAX	SPL SP ESP RSP	R8B R8W R8D R8	R12B R12W R12D R12
CL CX ECX RCX	BPL BP EBP RBP	R9B R9W R9D R9	R13B R13W R13D R13
DL DX EDX RDX	SIL SI ESI RSI	R10B R10W R10D R10	R14B R14W R14D R14
BL BX EBX RBX	DIL DI EDI RDI	R11B R11W R11D R11	R15B R15W R15D R15

FLAGS EFLAGS RFLAGS   IP EIP RIP
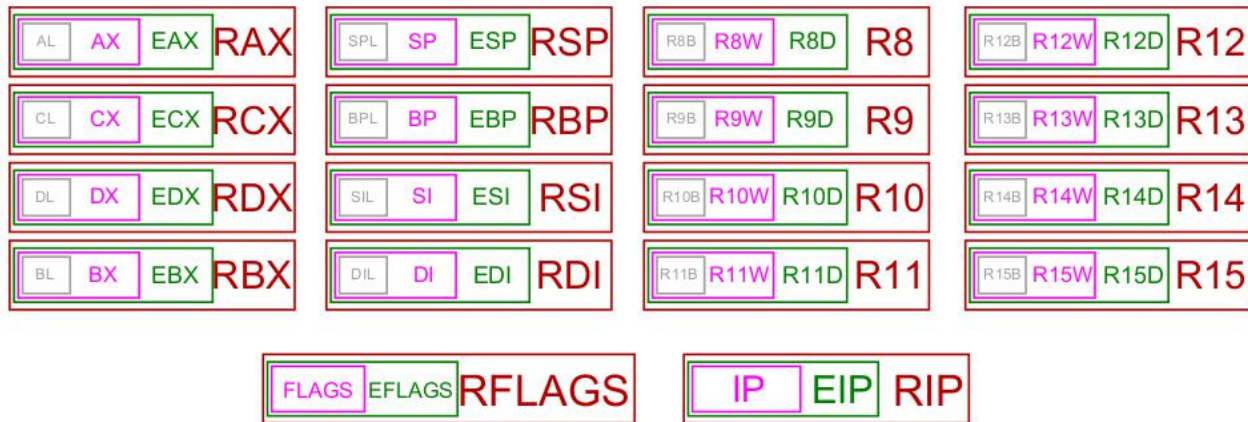
- 64-bit Register
- 32-bit Register
- 16-bit Register
- 8-bit Register

20 di 75 - Il Processore z64

x86_64 Registers

ZMM0	YMM0	XMM0	ZMM1	YMM1	XMM1
ZMM2	YMM2	XMM2	ZMM3	YMM3	XMM3
ZMM4	YMM4	XMM4	ZMM5	YMM5	XMM5
ZMM6	YMM6	XMM6	ZMM7	YMM7	XMM7
ZMM8	YMM8	XMM8	ZMM9	YMM9	XMM9
ZMM10	YMM10	XMM10	ZMM11	YMM11	XMM11
ZMM12	YMM12	XMM12	ZMM13	YMM13	XMM13
ZMM14	YMM14	XMM14	ZMM15	YMM15	XMM15

| ZMM16 | ZMM17 | ZMM18 | ZMM19 | ZMM20 | ZMM21 | ZMM22 | ZMM23 |
| ZMM24 | ZMM25 | ZMM26 | ZMM27 | ZMM28 | ZMM29 | ZMM30 | ZMM31 |

ST(0)	MM0	ST(1)	MM1
ST(2)	MM2	ST(3)	MM3
ST(4)	MM4	ST(5)	MM5
ST(6)	MM6	ST(7)	MM7

| CW | FP_IP | FP_DP | FP_CS |
| SW |
| TW |
| FP_DS |
| FP_OPC | FP_DP | FP_IP |

AL	AX	EAX	RAX	R8	R8W	R8D	R12	R12W	R12D	CR0	CR4
BL	BX	EBX	RBX	R9	R9W	R9D	R13	R13W	R13D	CR1	CR5
CL	CX	ECX	RCX	R10	R10W	R10D	R14	R14W	R14D	CR2	CR6
DL	DX	EDX	RDX	R11	R11W	R11D	R15	R15W	R15D	CR3	CR7
BP	EBP	RBP	DI	EDI	RDI	SP	ESP	RIP	CR3	CR8
SI	ESI	RSI	SP	ESP	RSP	MSW	CR9

CR10
CR11
CR12
CR13
CR14
CR15 MXCSR

| CS | SS | DS | GDTR | IDTR |
| ES | FS | GS | TR | LDTR |

RFLAGS EFLAGS FLAGS

DR0	DR6
DR1	DR7
DR2	DR8
DR3	DR9
DR4	DR10	DR12	DR14
DR5	DR11	DR13	DR15

Legend:
- 256-bit Register
- 80-bit Register
- 64-bit Register
- 16-bit Register
- 512-bit Register
- 128-bit Register
- 32-bit Register
- 8-bit Register

CR0

Bit	Name	Full Name	Description
0	PE	Protected Mode Enable	If 1, system is in protected mode, else system is in real mode
1	MP	Monitor co-processor	Controls interaction of WAIT/FWAIT instructions with TS flag in CR0
2	EM	Emulation	If set, no x87 FPU is present, if clear, x87 FPU is present
3	TS	Task switched	Allows saving x87 task context upon a task switch only after x87 instruction used
4	ET	Extension type	On the 386, it allowed to specify whether the external math coprocessor was an 80287 or 80387
5	NE	Numeric error	Enable internal x87 floating point error reporting when set, else enables PC style x87 error detection
16	WP	Write protect	When set, the CPU can't write to read-only pages when privilege level is 0
18	AM	Alignment mask	Alignment check enabled if AM set, AC flag (in EFLAGS register) set, and privilege level is 3
29	NW	Not-write through	Globally enables/disable write-through caching
30	CD	Cache disable	Globally enables/disable the memory cache
31	PG	Paging	If 1, enable paging and use the CR3 register, else disable paging

Entering Basic Protected Mode

- The code must set bit 0 (PE) of register CR0

- Setting PE to 1 does not immediately activate all its facilites

- It happens when the CS register is first updated

- This can be only done using a far jump (`ljmp`) instruction, as already mentioned.

- After this, code executes in 32/64-bit mode

Entering Basic Protected Mode

```
 ljmp 0x0000, PE_mode


 .code32
PE_mode:
 # Set up the protected-mode data segment registers
 movw $PROT_MODE_DSEG, %ax
 movw %ax, %ds
 movw %ax, %es
 movw %ax, %fs
 movw %ax, %gs
 movw %ax, %ss
```

Segment Registers in Protected Mode

- In Protected Mode, a segment is no longer a raw number

- It contains (also) an index into a table of **segment descriptors**

- There are three types of segments:
  - code
  - data
  - system

Descriptor Table Entry

63		55		47			32
4 bytes	Base (24-31)	G		Limit (16-19)	1	DPL	Base (16-23)
4 bytes	Base (0-15)			Limit (0-15)			
31				15			0
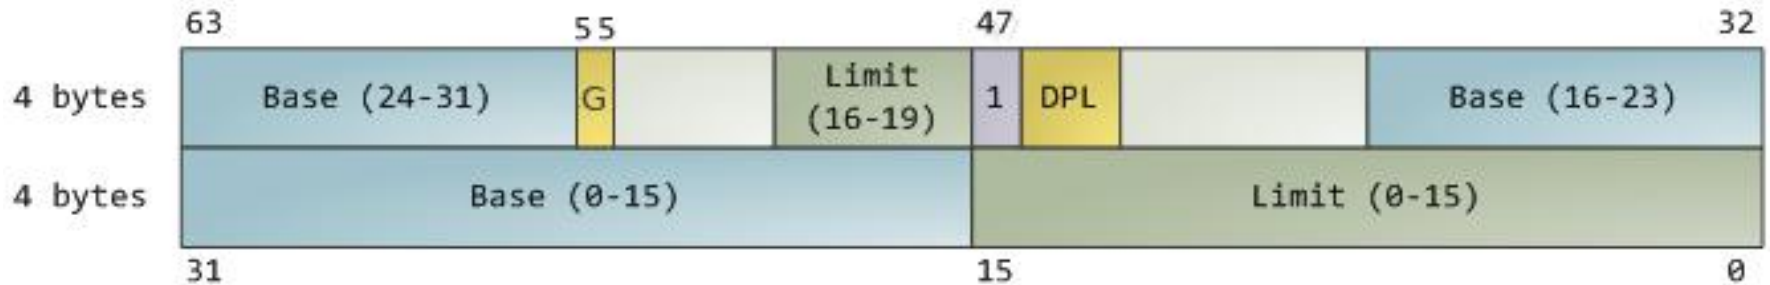
- **Base**: 32-bit linear addressing pointing to the beginning of the segment

- **Limit**: size of the segment

- **G**: *Granularity*. If set, size is to be multiplied by 4096

- **Descriptor Privilege Level** (DPL): a number in [0-3] to control access to the segment

Protected Mode: Privilege Levels
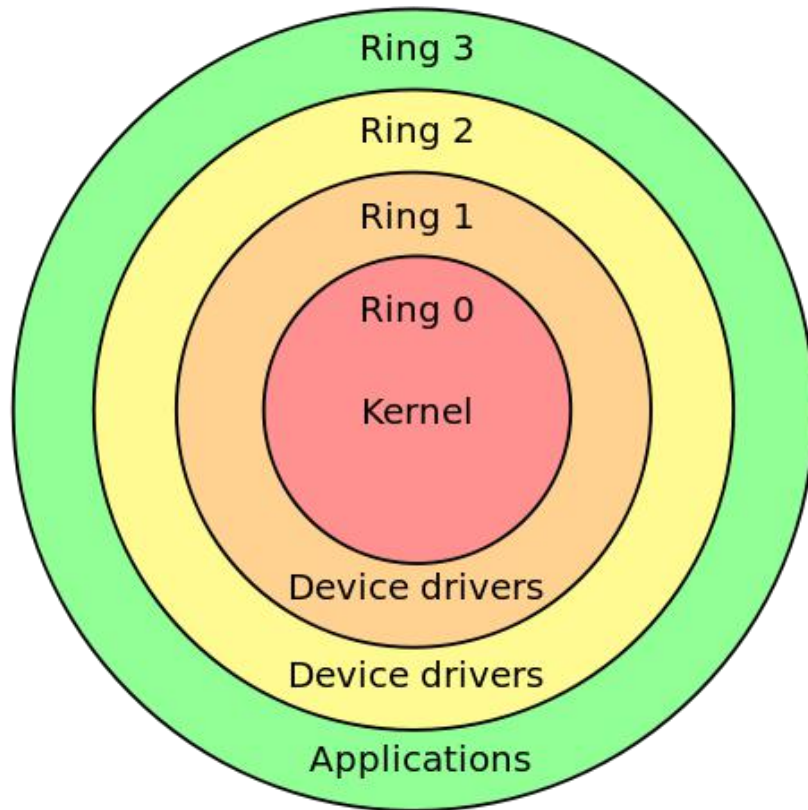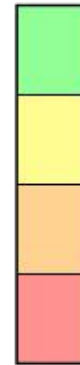
Ring 3 has restricted access to memory management, instructions execution (around 15 allowed only at ring 0), and I/O ports

Descriptor Tables

- Two tables are available on x86 architectures

- **Global Descriptor Table** (GDT):
  - This is a system-wide table of descriptors
  - It is pointed by the GDTR register

- **Local Descriptor Table** (LDT):
  - Pointed by the LDTR register
  - Not used anymore

Segment Selectors

- **TI**: set to 0 for the GDT, set to 1 for the LDT
- **Index**: specifies the segment selector within the associated table
- **Requested Privilege Level** (RPL): we'll come to that later

Segmented Addressing Resolution

Segmented Addressing Resolution

Segmentation cannot be disabled
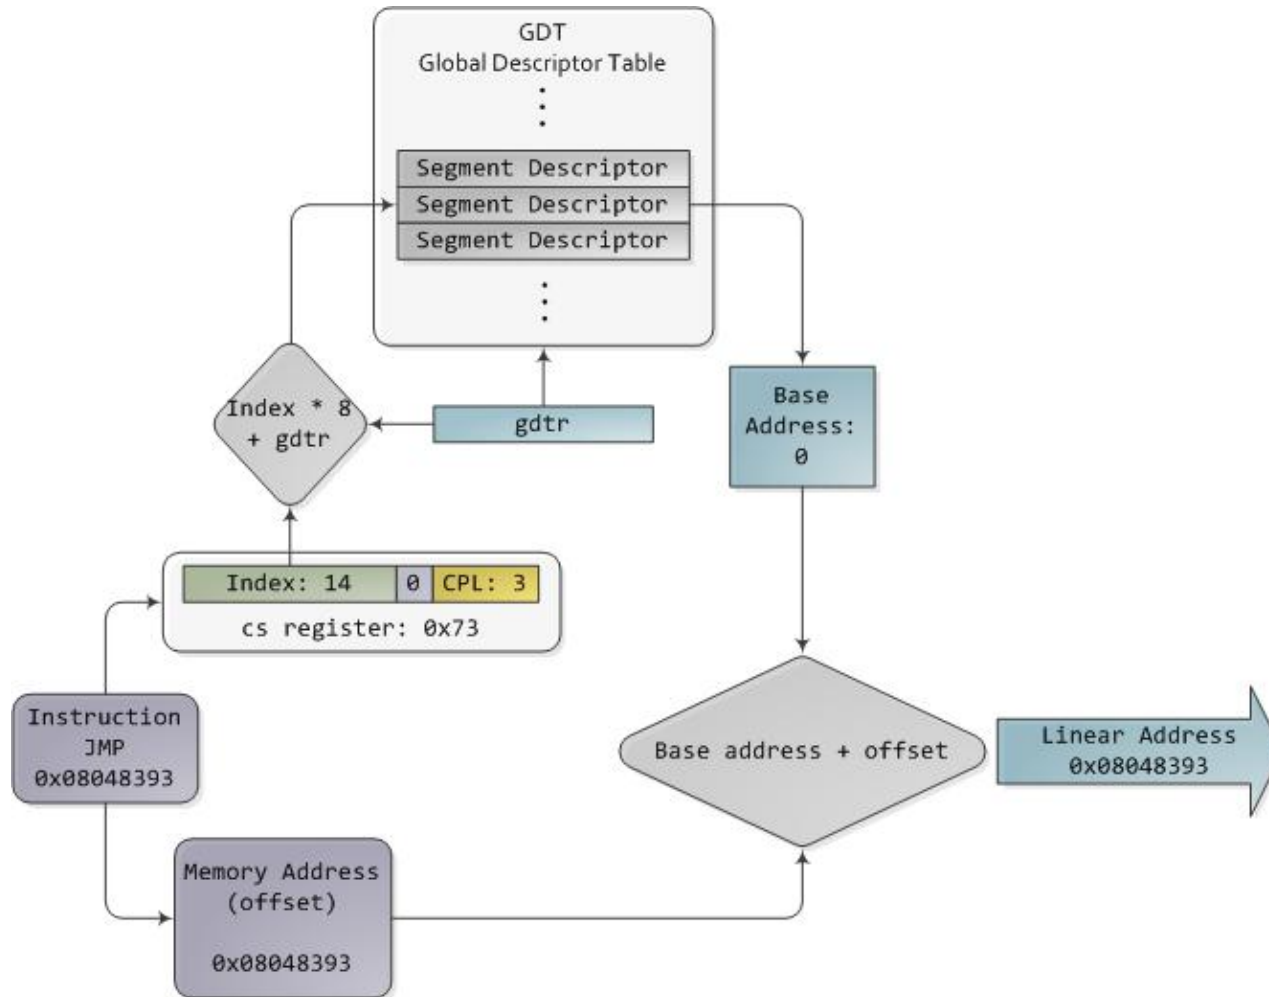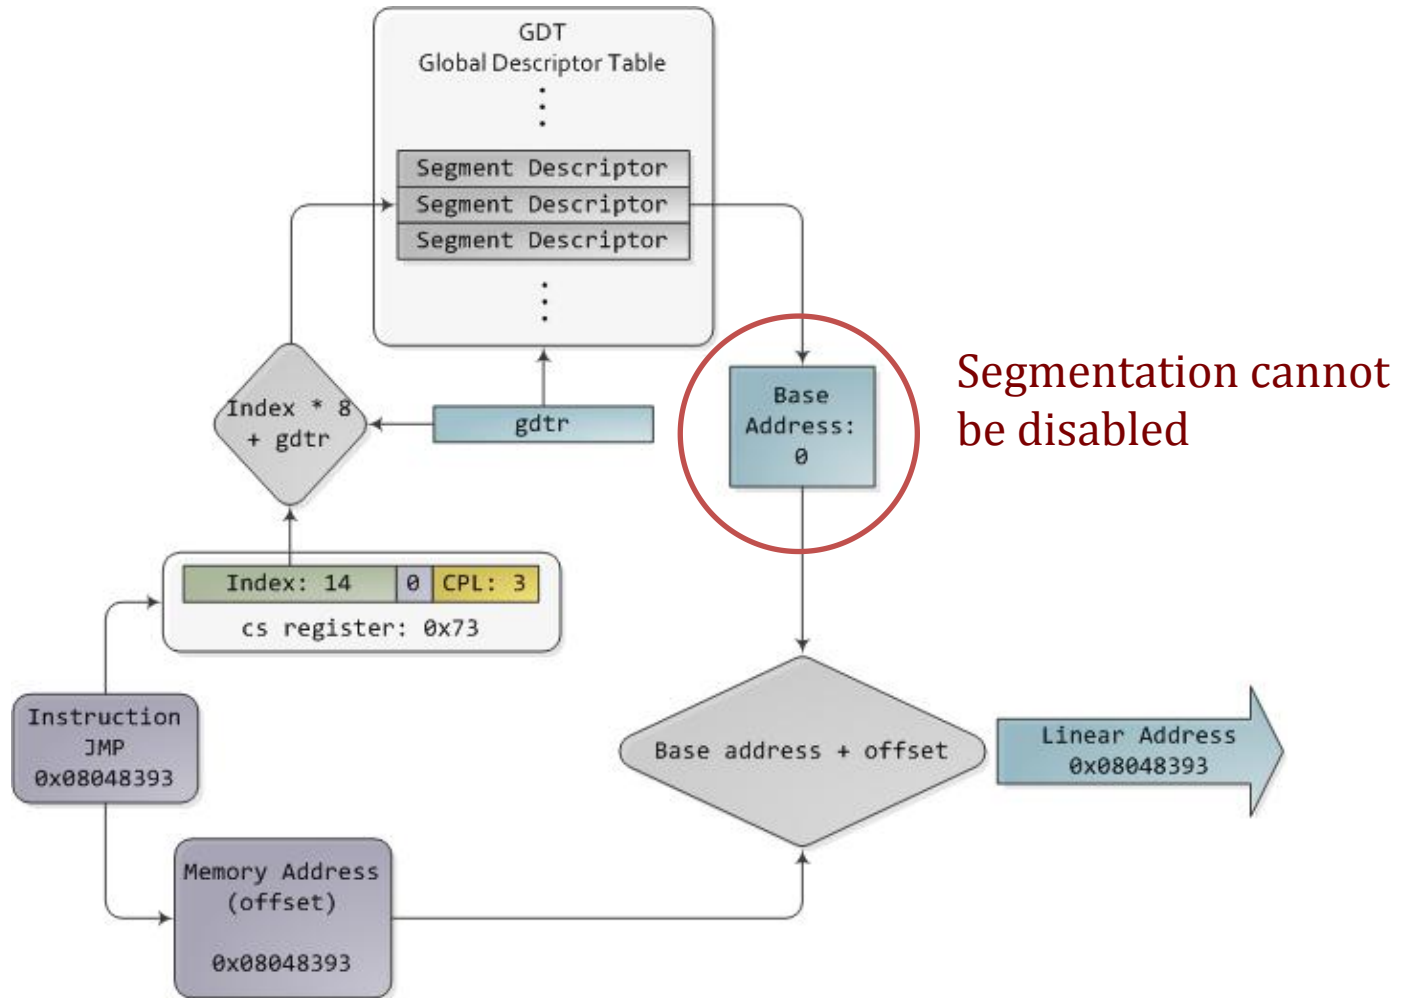
Segment Caching

- Accessing the GDT for every memory access is not performance-wise

- Segment registers have a non-programmable hidden part to store the cached descriptor
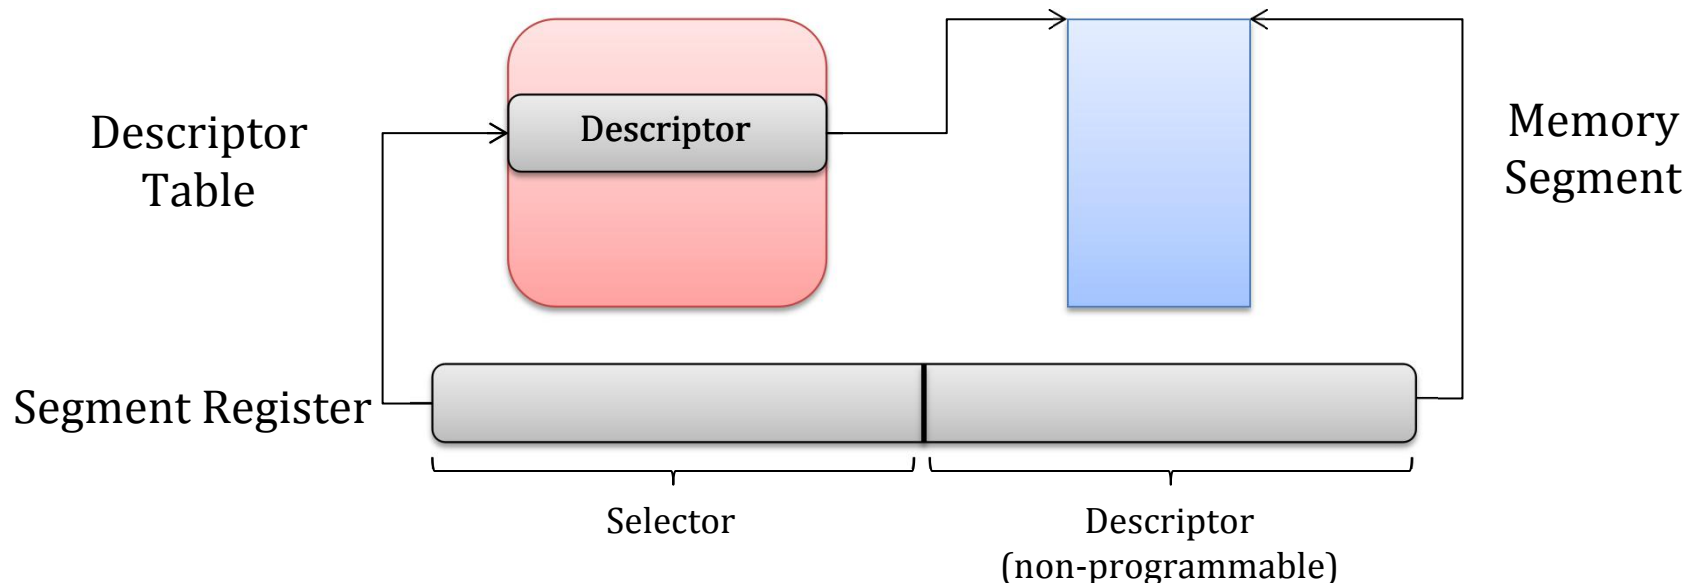
x86 Enforcing Protection

- A Descriptor Entry has a DPL

- The firmware must check if an access to a certain segment is allowed

- There must be a way to change current privilege

Ring 3
Ring 2
Ring 1
Ring 0
Kernel
Device drivers
Device drivers
Applications

Change always allowed

Should be controlled/denied
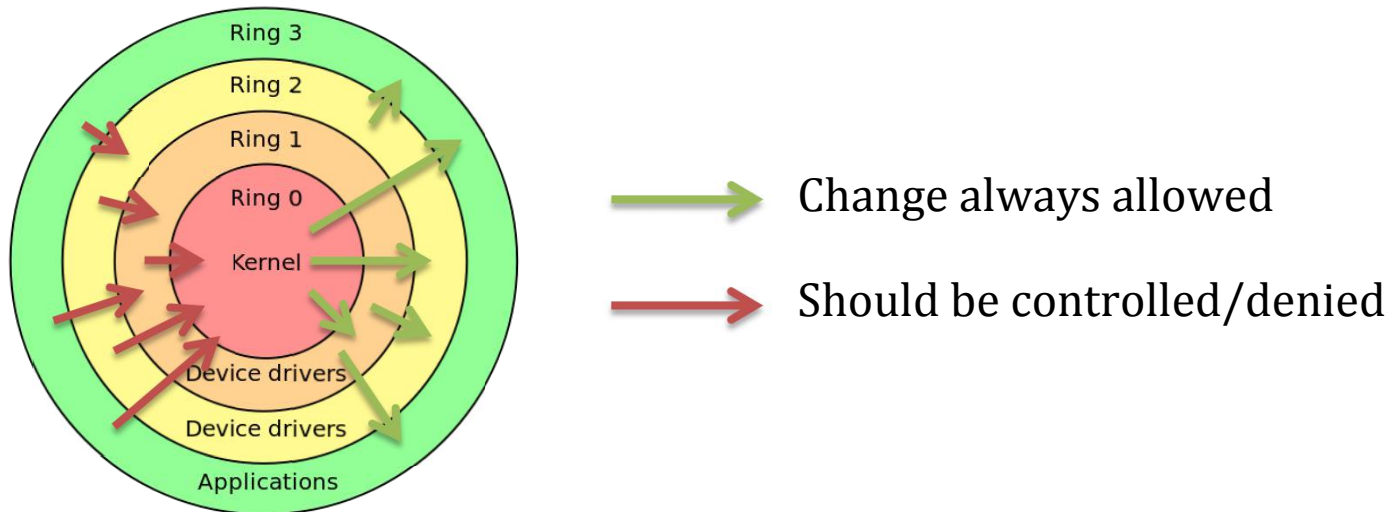
Data Segment vs Code Segment

- RPL is present only in data segment selectors (e.g. SS or DS)

- **Current Privilege Level** (CPL): this is only in CS, which can be loaded only with a `ljmp`/`lcall`

- Overall we have 3 different privilege-level fields:
  CPL, RPL, and DPL

Data segment selector

16 bits | Index (3-15) | T I | RPL
15 | | 2 | 0

Code segment selector

16 bits | Index (3-15) | T I | CPL
15 | | 2 | 0

Protection upon Segment Load

- CPL is managed by the CPU: it's *always* equals to the current CPU privilege level

- CPU Memory protection comes at two points:
  - Memory access via a linear address
  - Data segment selector load operation

Protection upon Segment Load

Current code segment register

CPL

Data segment selector being loaded

Index | T I | RPL

Selects

Segment descriptor

DPL

For SS, the condition is:
CPL = RPL = DPL

MAX(CPL, RPL) <= DPL

True:
Segment load ok

False:
General-protection exception
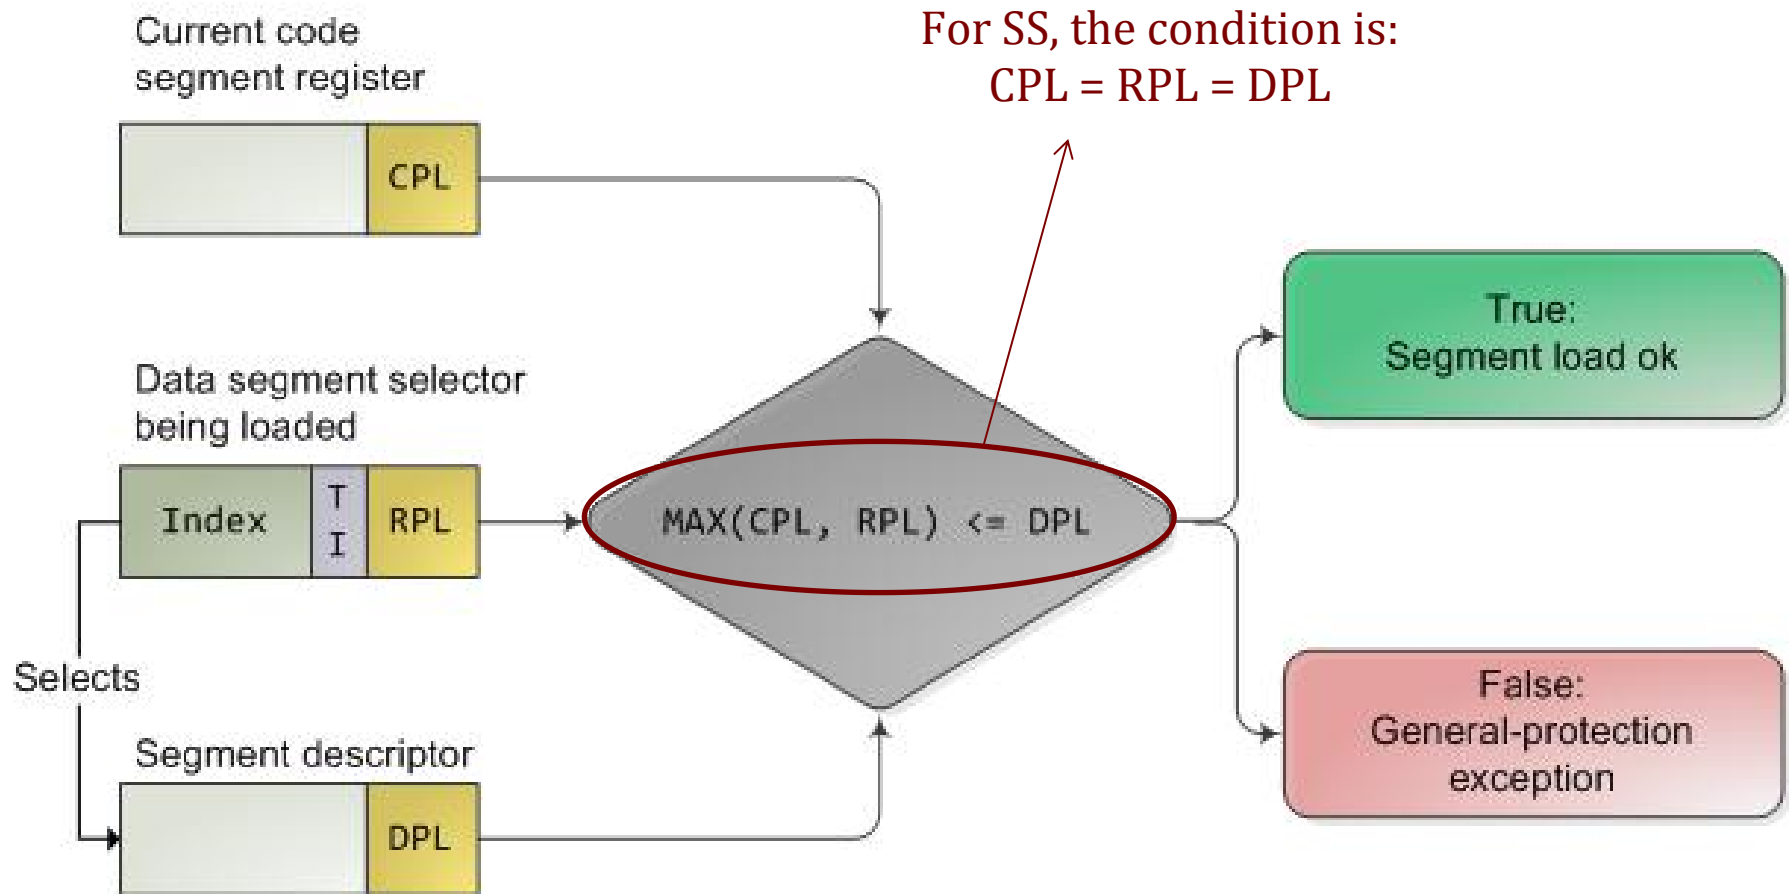
Getting Higher Privileges

- Accessing segment with a higher privilege (lower ring) with not control might allow malicious code to subvert the kernel

- To control transfer, code must pass through a controlled **gate**

- **Gate descriptors** are used to identify possible gates through which control can pass

Controlled Access Through Gates
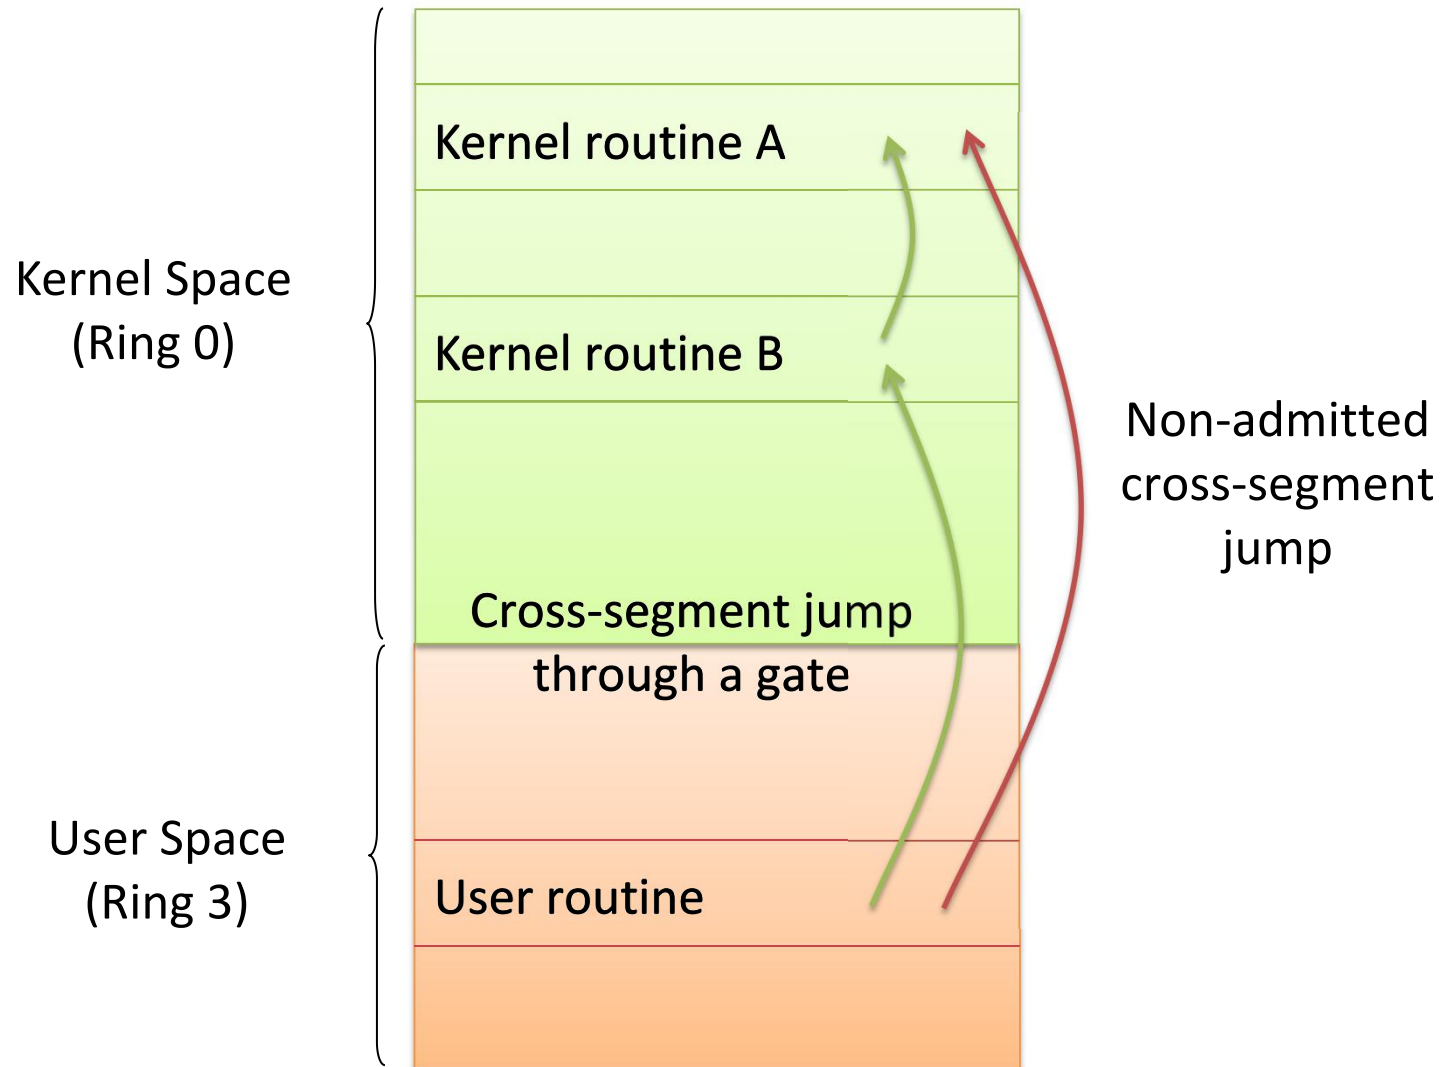
Gate Descriptors

- A gate descriptor is a segment descriptor of type *system*:
    - Call-gate descriptors
    - **Interrupt-gate descriptors**
    - **Trap-gate descriptors**
    - Task-gate descriptors

- These are referenced by the **Interrupt Descriptor Table** (IDT), pointed by the IDTR register

Interrupts vs Traps

- Interrupts are **asynchronous events** not related to the current CPU execution flow

- Interrupts are generated by external devices, and can be masked or not (NMI)

- Traps (or *exceptions*) are **synchronous events**, strictly related to the current CPU execution (e.g. division by zero)

- Traps are the historical way to demand access to kernel mode

IDT and GDT Relations
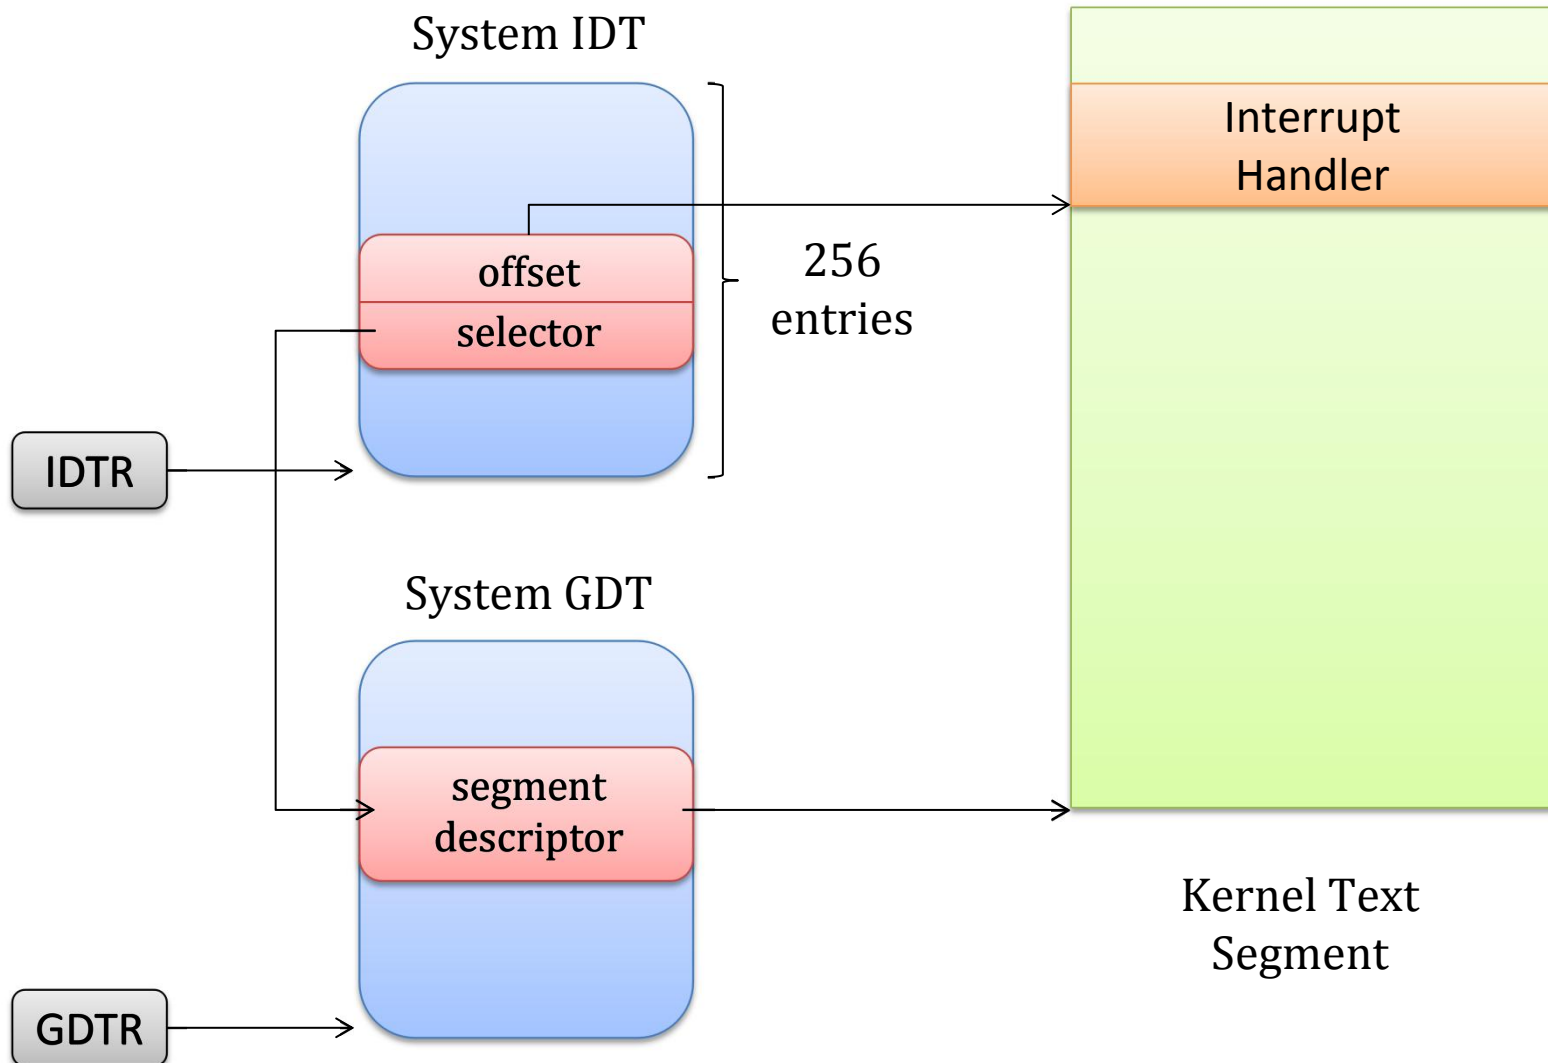
GDT in Linux

Linux's GDT	Segment Selectors
null	0x0
reserved	
reserved	
reserved	
not used	
not used	
TLS #1	0x33
TLS #2	0x3b
TLS #3	0x43
reserved	
reserved	
reserved	
kernel code	0x60 (__KERNEL_CS)
kernel data	0x68 (__KERNEL_DS)
user code	0x73 (__USER_CS)
user data	0x7b (__USER_DS)

Linux's GDT	Segment Selectors	
TSS	0x80	← Different for all cores
LDT	0x88	← Shared across all cores
PNPBIOS 32-bit code	0x90	
PNPBIOS 16-bit code	0x98	
PNPBIOS 16-bit data	0xa0	
PNPBIOS 16-bit data	0xa8	
PNPBIOS 16-bit data	0xb0	
APMBIOS 32-bit code	0xb8	
APMBIOS 16-bit code	0xc0	
APMBIOS data	0xc8	
not used		
not used		
not used		
not used		
not used		
double fault TSS	0xf8	

There is one copy of this table for each core

Task State Segment (TSS)

- Its an x86 structure holding information about a task

- It is intended to handle task management

- It stores:
  - Processor registers state
  - I/O Port Permissions
  - Inner-level Stack Pointers
  - Previous TSS link

Task State Segment (TSS)

- It can be everywhere in memory (hence the GDT entry required to access it)

- On Linux, it's in kernel data memory

- Each TSS is stored in the `int_tss` array.

- The selector is kept in the Task Register (TR)

- It can be loaded using the privileged `ltr` instruction

Task State Segment (TSS)

Pointer to a bitmap →

31		15		0	
I/O Map Base Address				T	100
		LDT Segment Selector			96
		GS			92
		FS			88
		DS			84
		SS			80
		CS			76
		ES			72
EDI					68
ESI					64
EBP					60
ESP					56
EBX					52
EDX					48
ECX					44
EAX					40
EFLAGS					36
EIP					32
CR3 (PDBR)					28

CPU State

		SS2			24
ESP2					20
		SS1			16
ESP1					12
		SS0			8
ESP0					4
		Previous Task Link			0

Privilege-level stacks

Reserved bits. Set to 0.
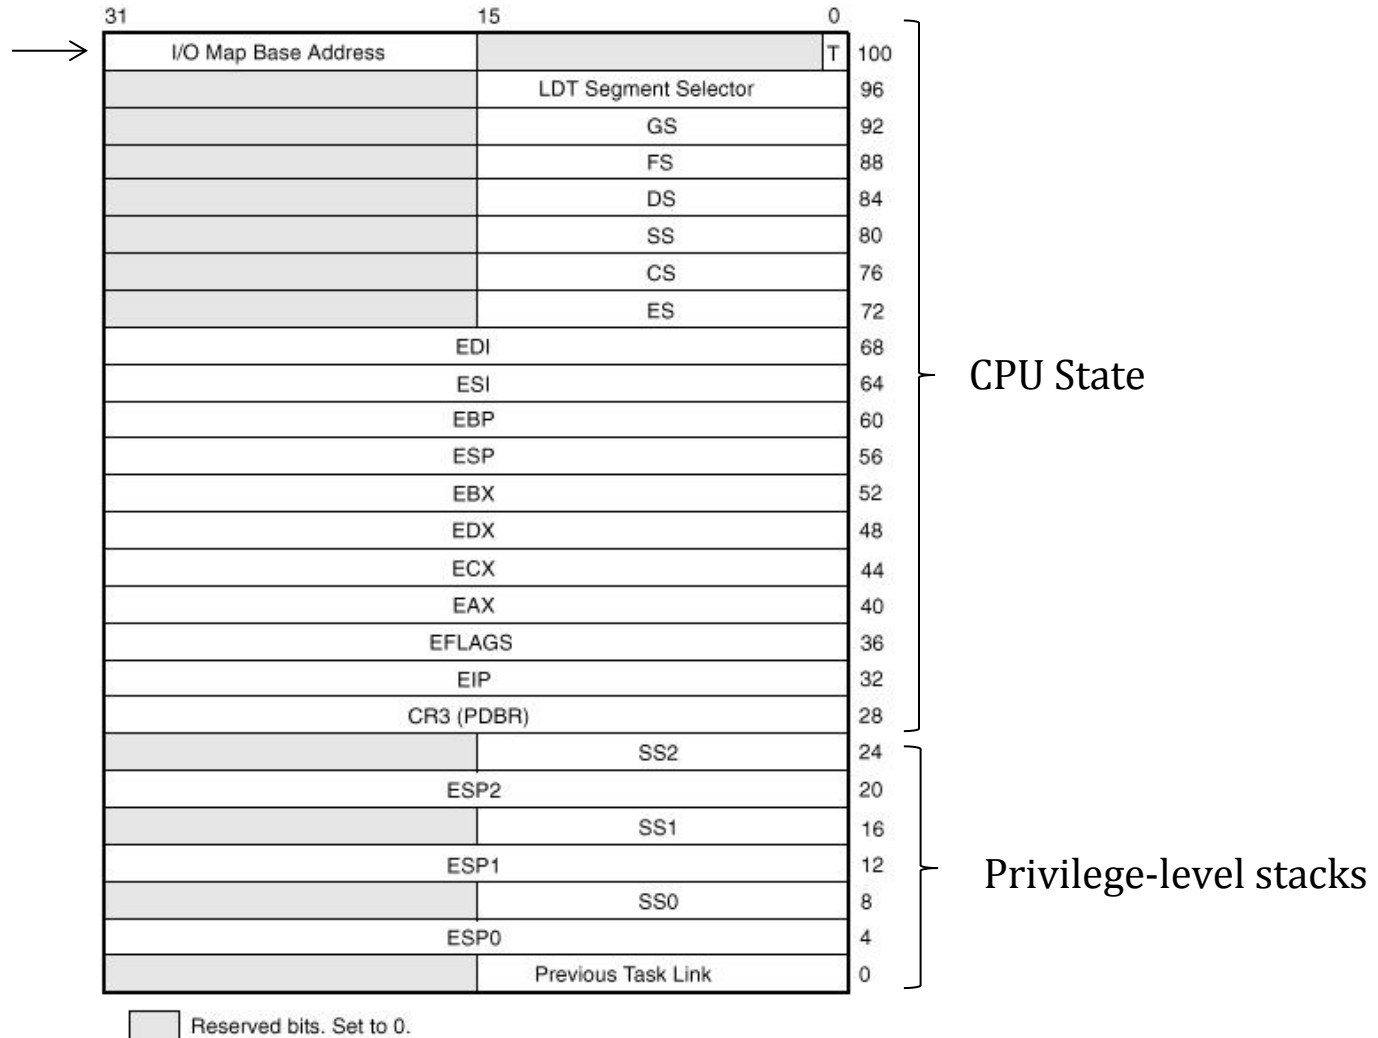
Task State Segment (TSS)

- The *Base* field within the *n*-th core TSS register points to the *n*-th entry of the `int_tss` array

- *G*=0 and *Limit*=0xeb

  – given that TSS is 236 bytes in size

- *DPL*=0

  •TSS cannot be accessed in user mode

Entering Ring 0 from Ring 3

Current code segment register
(__USER_CS or __KERNEL_CS in Linux)

CPL: 3
or 0

Destination Code Segment
(__KERNEL_CS in Linux)

DPL: 0

CPL >= Destination segment DPL

For software-generated interrupt,
CPL <= gate DPL

Selects

63   47   32

Offset (31-16)   1   DPL   Gate DPL is 3 or 0

Segment Selector: 0x60 (index is 12)   Offset (15-0)

31   15   0

Interrupt-gate/trap-gate descriptor
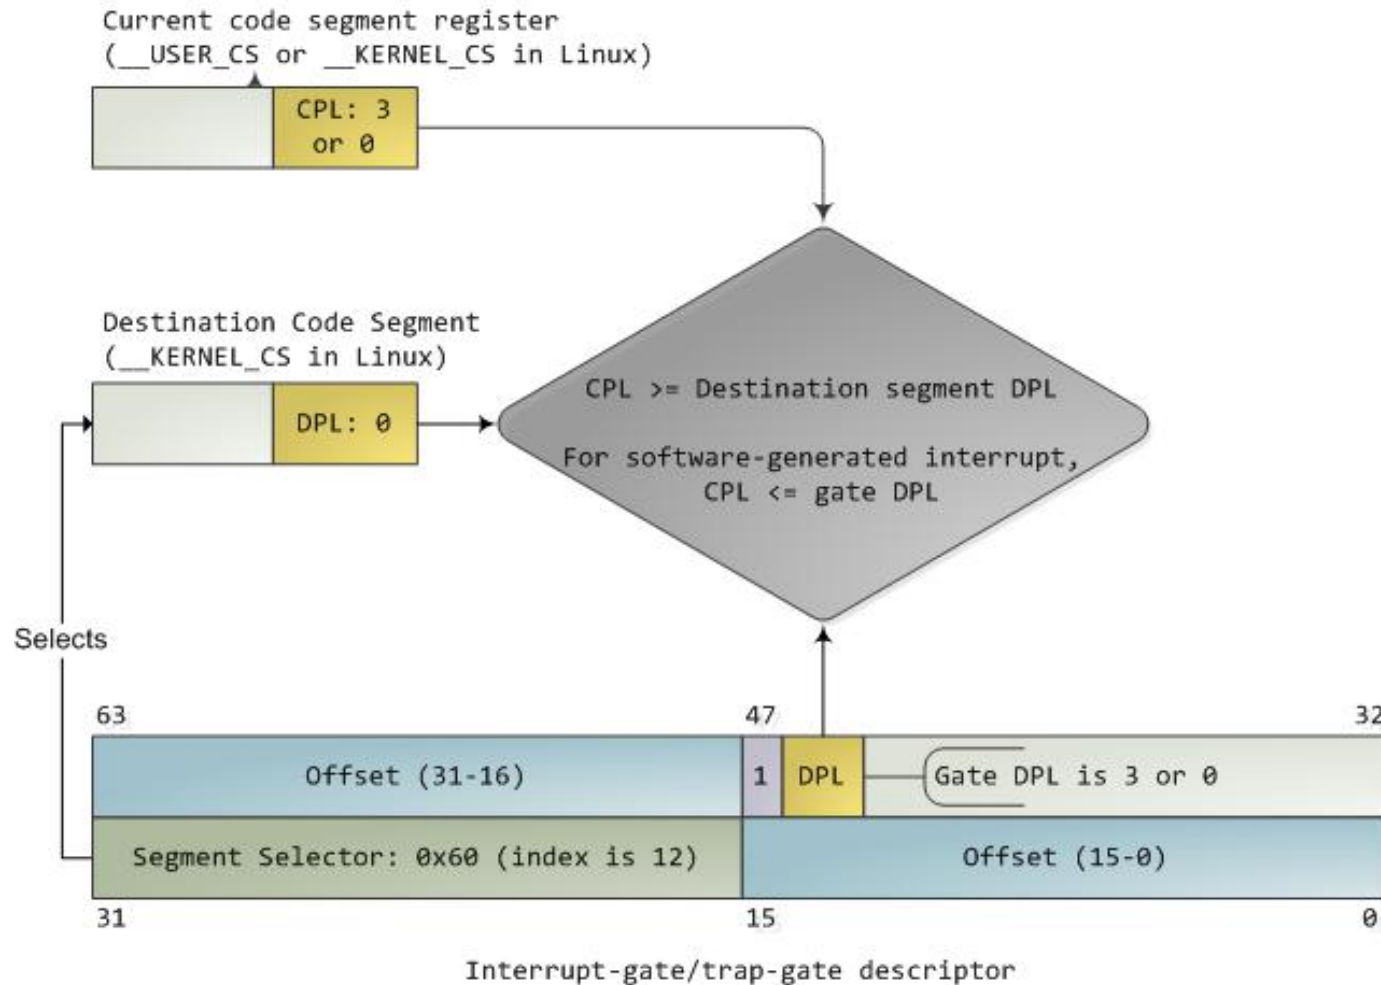
Protected Mode Paging

- Since 80386, x86 CPUs add an additional step in address translation

Memory Address Translation
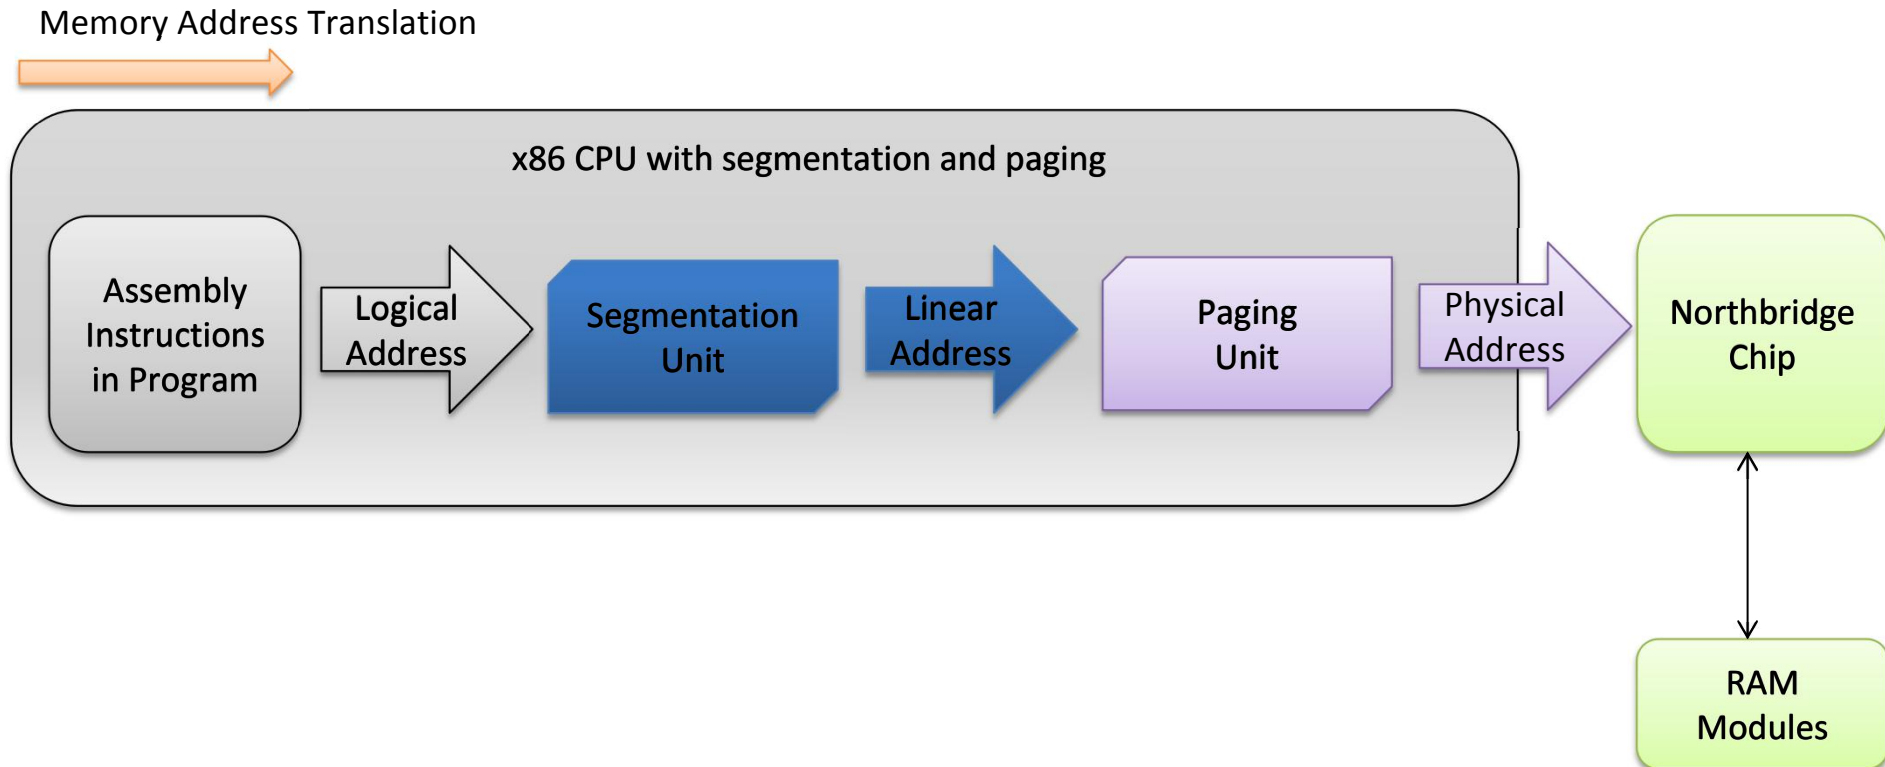
Protected Mode Paging

- Paging has to be explicitly enabled
  - Entering Protected Mode does not enable it automatically
  - Several data structures must be setup before

- Paging allows to manage memory protection at a smaller granularity than segmentation

i386 Paging Scheme

i386 Paging Scheme

- Both levels entail 4 KB memory blocks

- Each block is an array of 4-byte entries

- Hence we can map **1 K x 1K pages**

- Since each page is 4 KB in sixe, we get a 4 GB virtual addressing space

i386 PDE entries

Page-Directory Entry (4-KByte Page Table)

31	12	11	9	8	7	6	5	4	3	2	1	0
Page-Table Base Address		Avail		G	P S	0	A	P C D	P W T	U / S	R / W	P

Available for system programmer's use
Global page (Ignored)
Page size (0 indicates 4 KBytes)
Reserved (set to 0)
Accessed (Sticky bit)
Cache disabled
Write-through
User/Supervisor
Read/Write
Present

i386 PTE entries

Page-Table Entry (4-KByte Page)

31	12 11	9 8	7	6	5	4	3	2	1	0
Page Base Address	Avail	G	P A T	D	A	P C D	P W T	U / S	R / W	P

Available for system programmer's use
Global Page (TLB caching policy)
Page Table Attribute Index
Dirty
Accessed (Sticky bit)
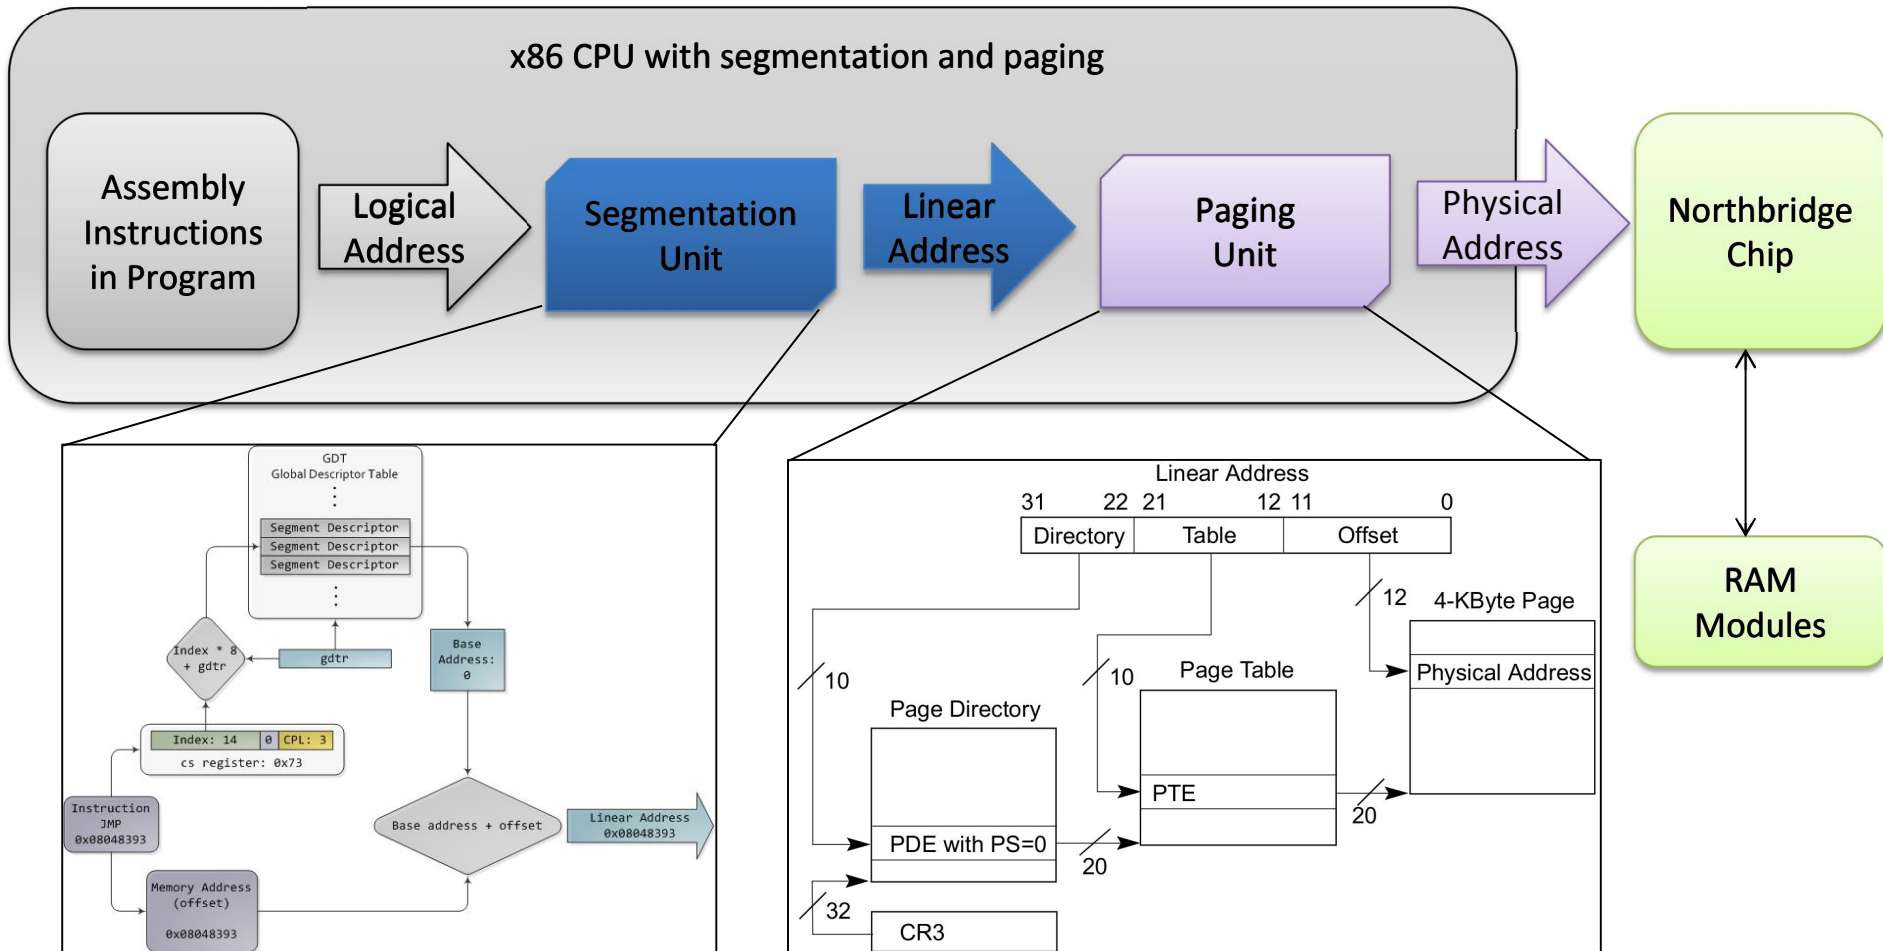Cache Disabled
Write-Through
User/Supervisor
Read/Write
Present

Translation Lookaside Buffer

Translation Lookaside Buffer

Relations to Trap/Interrupt Table

- Upon a TLB miss, firmware accesses the page table

- The first checked bit is PRESENT

- If this bit is zero, a page fault occurs which gives rise to a trap

- CPU registers (including EIP and CS) are saved on the system stack

- They will be restored when returning from trap: the trapped instruction is re-executed

- **Re-execution might give rise to additional traps, depending on firmware checks on the page table**

- As an example, the attempt to access a read only page in write mode will give rise to a trap (which triggers the **segmentation fault handler**)

Linux memory layout on i386

0xFFFFFFFF ... 0xC0000000	Kernel (1 GB)
0xBFFFFFFF ... 0x00000000	

User Space
(3 GB)
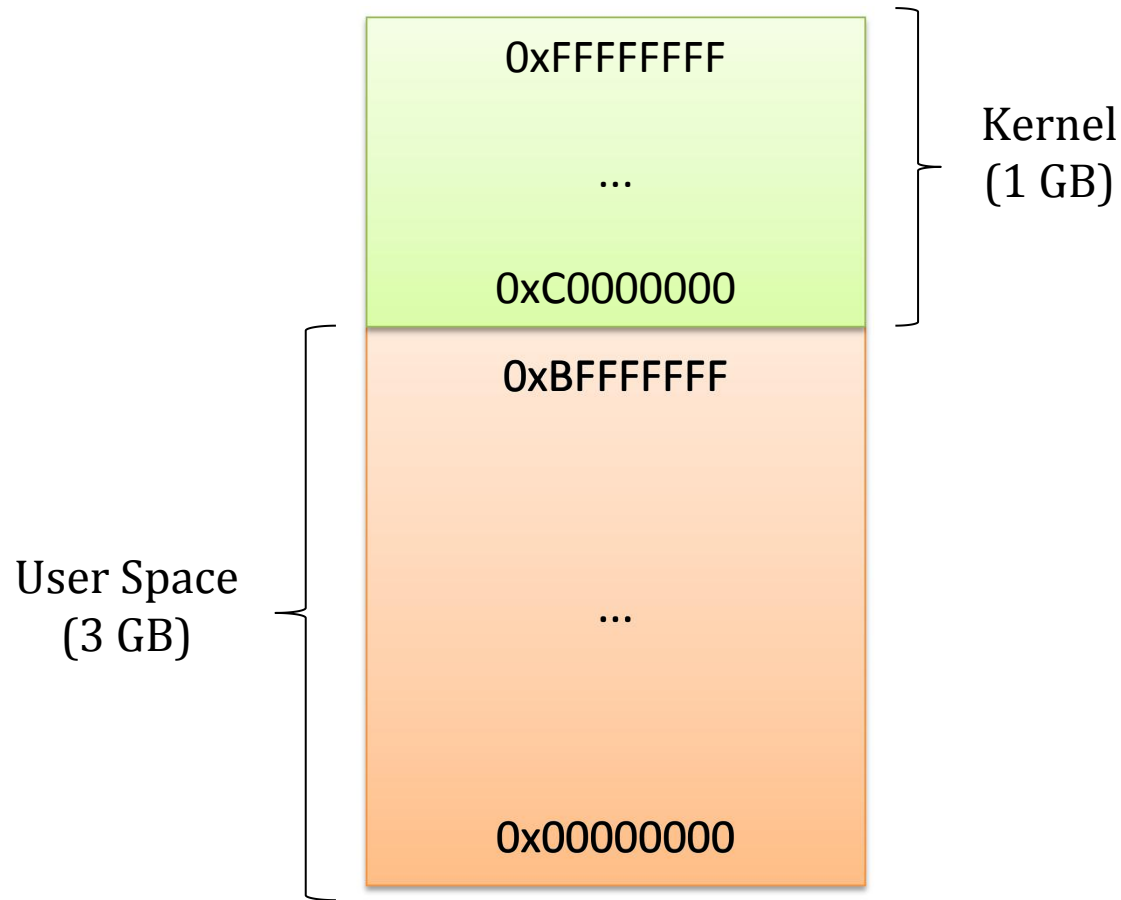
Physical Address Extension (PAE)

- An attempt to extend over the 4GB limit on i386 systems

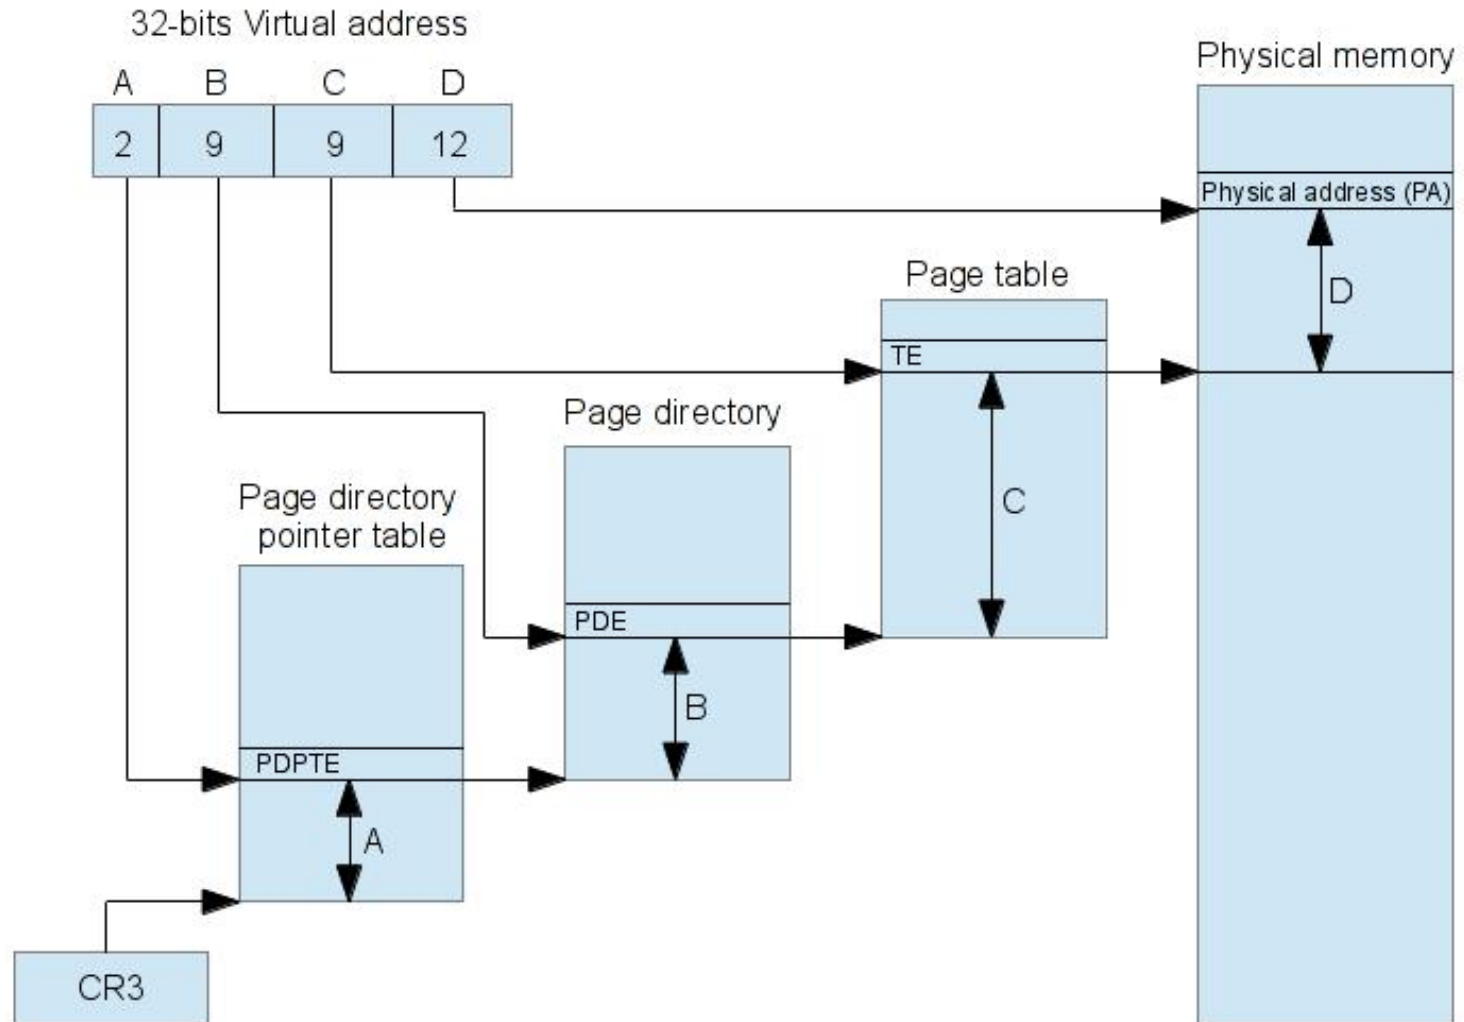- Present since the Intel Pentium Pro

- Supported on Linux since kernel 2.6

- Addressing is extended to 36 bits

- This allows to drive up to 64 GB of RAM memory

- Paging uses 3 levels

- CR4.PAE-bit (bit 5) tells if PAE is enabled

Physical Address Extension (PAE)

x64 Paging Scheme

- PAE is extended via the so called "long addressing"

- 2^{64} bytes of logical memory in theory

- Bits [49-63] are short-circuited to 1
  - Up to 2^{48} canonical form addresses (lower and upper half)
  - A total of 256 TB addressable memory

- Linux currently allows for 128 TB of logical addressing of individual processes and 64 TB for physical addressing

Canonical Addresses

64-bit

48-bit

Linux memory layout on x64

TEXT

DATA

Kernel

Stack
Heap
Shared Objects

0xFFFF FFFF FFFF FFFF

0xFFFF 8000 0000 0000

0x0000 7FFF FFFF FFFF

Non-canonical
addresses

0x0000 0000 0000 0000
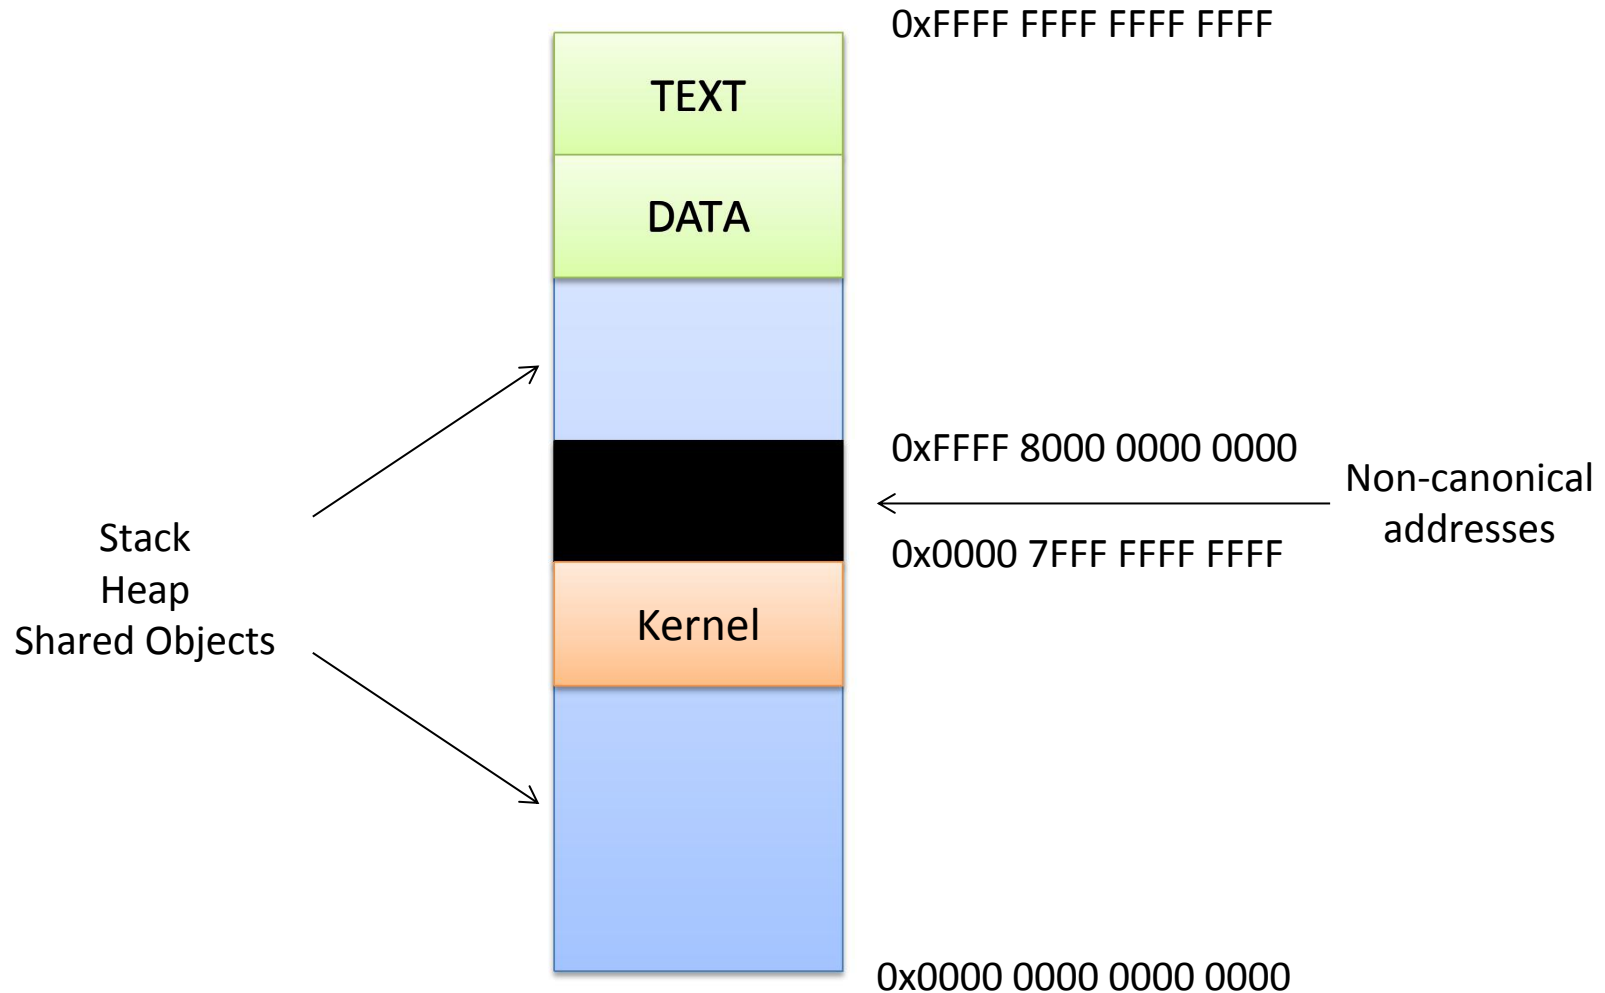
48-bit Page Table (4KB pages)

Also referred to as Page General Directory (PGD)
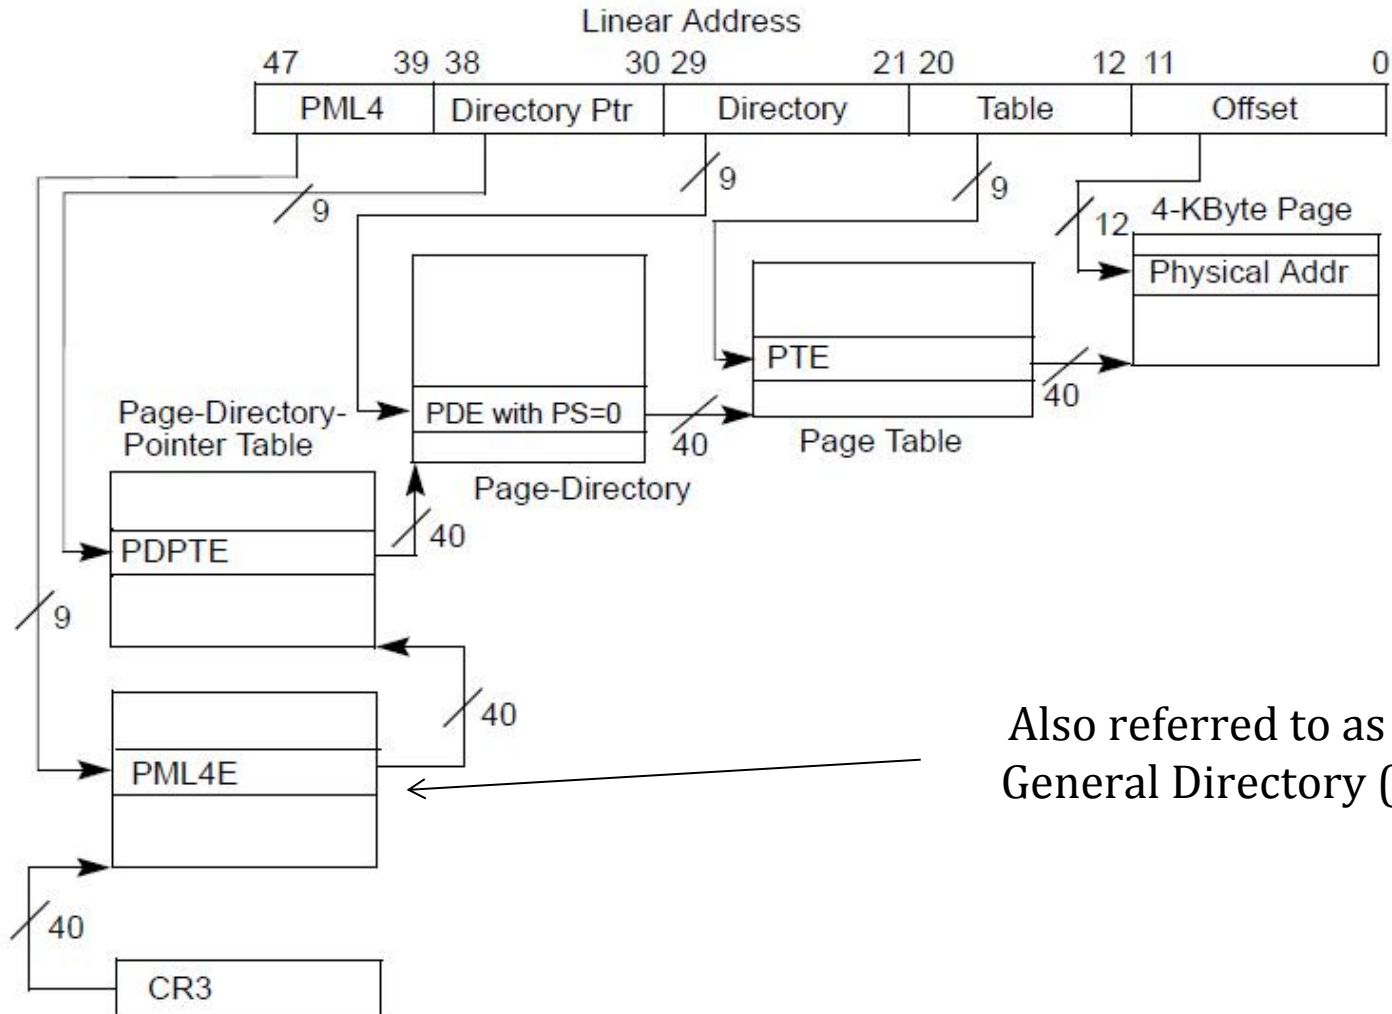
CR3 and Paging Structure Entries

Huge Pages

- Ideally x64 processors support them starting from PDPT

- Linux typically offers the support for huge pages pointed by the PDE (page size 512*4KB)

- See: `/proc/meminfo` and `/proc/sys/vm/nr_hugepages`

- These can be mmap'ed via file descriptors and/or `mmap` parameters (e.g. `MAP_HUGETLB` flag)

- They can also be requested via the `madvise(void *, size_t, int)` system call (with `MADV_HUGEPAGE` flag)

Back to the bootstrap

- We have enabled Protected Mode

- Addressing mode is still using segmentation

- We can address up to 16 Mb of RAM

- Paging is not yet setup

- It's time to start loading the actual kernel

Boot Sequence

BIOS/UEFI	The actual Hardware Startup
Bootloader Stage 1	Executes the Stage 2 bootloader (skipped in case of UEFI)
Bootloader Stage 2	Loads and starts the Kernel
Kernel Startup	The Kernel takes control of and initializes the machine (machine−dependent operations)
Init	First process: basic environment initialization (e.g., SystemV Init, systemd)
Runlevels/Targets	Initializes the user environment (e.g., single−user mode, multiuser, graphical, ...)

Second Stage Bootloader

- There are various versions of this software
  - In GRUB it is GRUB Stage 2
  - In Win NT it is `c:\ntldr`
- The second stage bootloader reads a configuration file, e.g. to startup a boot selection menu
  - `grub.conf` in GRUB, `boot.ini` in Win NT
- The kernel initial image is loaded in memory using BIOS disk I/O services
  - For Linux, it is `/boot/vmlinuz-*`
  - For Win NT, it is `c:\Windows\System32\ntoskrnl.exe`

Historical Linux Bootcode

- The historical bootsector code for LINUX (i386) is in `arch/i386/bootsect.S` (no longer used)

- It loaded `arch/i386/bootsetup.S` and the kernel image in memory

- The code in `arch/i386/bootsetup.S` initialized the architecture (e.g. the CPU state for the actual kernel boot)

- It ultimately gave control to the initial kernel image

Unified Extensible Firmware Interface (UEFI)

- Modular (you can extend it with drivers)

- Runs on various platforms

- It's written in C

- It supports a bytecode (portability to other architectures)

- It's completely different from BIOS

UEFI Boot

- UEFI boot manager takes control right after the system is powered on
- It looks at the boot configuration
- It loads the firmware settings into RAM from nvRAM
- Startup files are stored on a dedicated EFI System Partition (ESP)
  - It'a a FAT32 partition
  - It has one folder for each OS on the system
- MBR cannot handle disks larger than 2TB

UEFI Boot

- It can automatically detect new uefi-boot targets
  - UEFI uses standard path names
    - `/efi/boot/boot_x64.efi`
    - `/efi/boot/bootaa64.efi`

- UEFI programs can be easily written

```
#include <efi.h>
#include <efilib.h>

EFI_STATUS EFIAPI
efi_main(EFI_HANDLE ImageHandle, EFI_SYSTEM_TABLE *SystemTable) {
        InitializeLib(ImageHandle, SystemTable);
        Print(L"Hello World\n");

        return EFI_SUCCESS;
}
```

GUID Partition Table
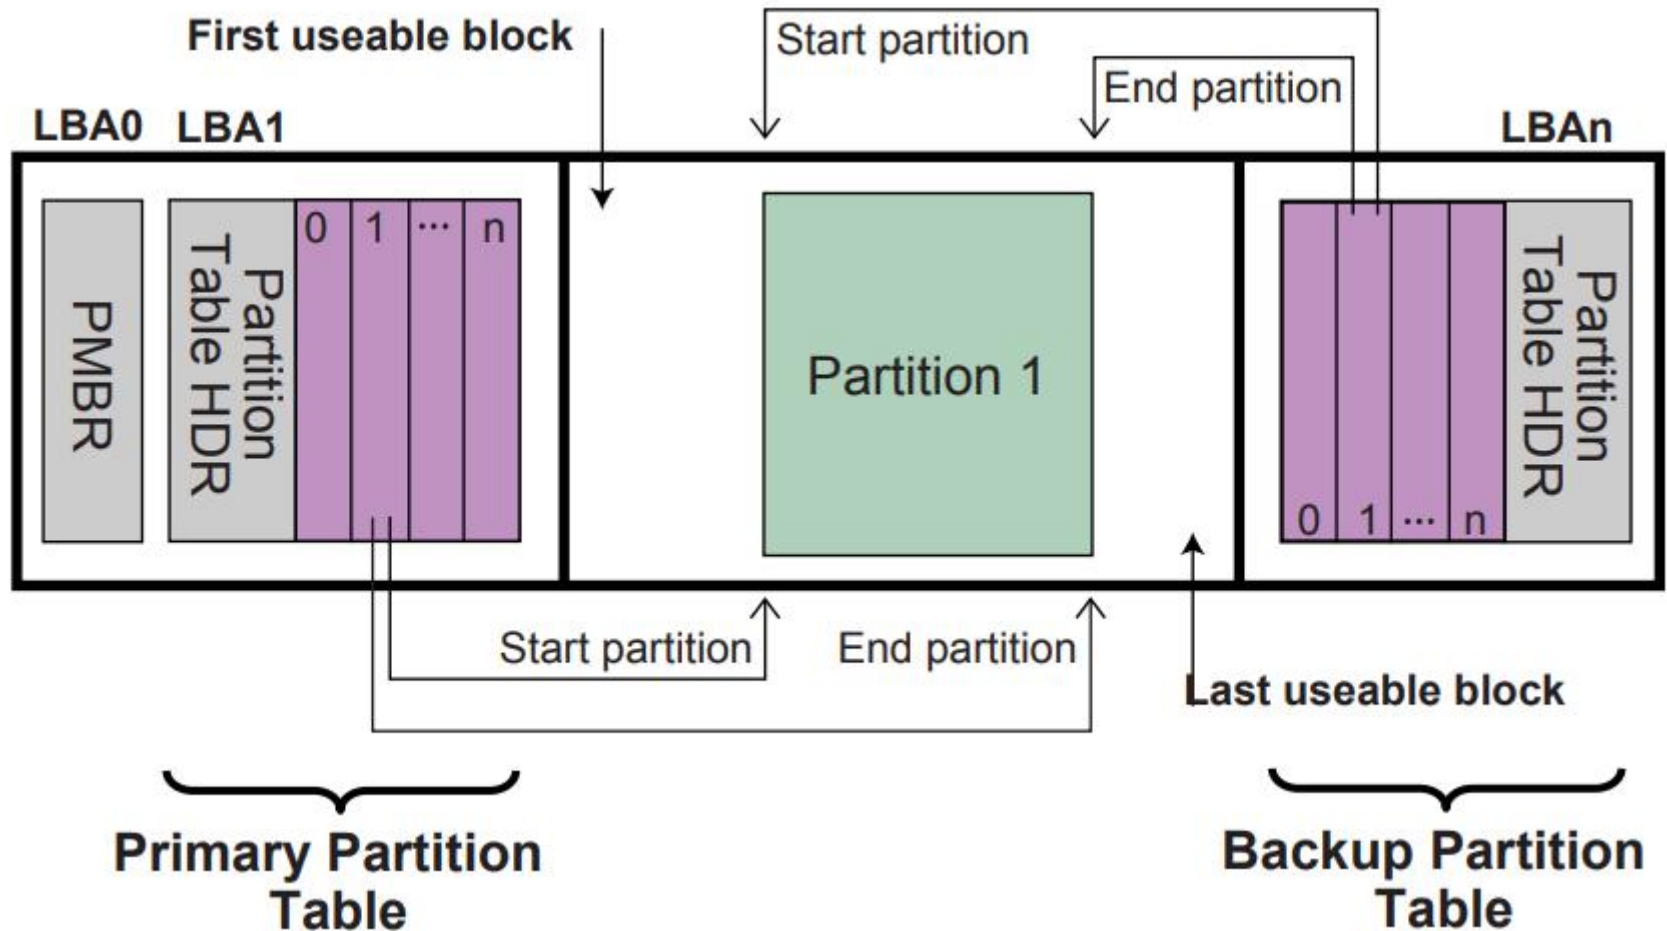
Secure Boot

- There is a kind of malware which takes control of the system before the OS starts
  – MBR RootKits
- Usually, these RootKits hijack the IDT for I/O operations, to execute their own wrapper
- When the kernel is being loaded, the RootKit notices that and patches the binary code while loading it into RAM

Secure Boot

- UEFI allows to load only signed executables
- Keys to verify signatures are installed in UEFI configuration
  - Platform Keys (PK): tells who "owns and controls" the hardware platform
  - Key-Exchange Keys (KEK): shows who is allowed to update the hardware platform
  - Signature Database Keys (DB): show who is allowed to boot the platform in secure mode

Dealing with multicores

- Who shall execute the startup code?
- For legacy reasons, the code is purely sequential
- Only one CPU core (the master) should run the code

- At startup, only one core is active, the others are in an idle state
- The startup procedure has to wake up other cores before starting the actual kernel

Interrupts on Multicore Architectures

- The Advanced Programmable Interrupt Controller (APIC) is used for sophisticated interrupt sending/redirection

- Each core has a Local APIC (LAPIC) controller, which can send Inter-Processor Interrupts (IPIs)
  - LAPICs are connected through the (logical) "APIC Bus"
  - LINT 0 : normal interrupts — LINT 1 : Non-maskable Interrupts

- I/O APICs contain a redirection table, which is used to route the interrupts it receives from peripheral buses to one or more local APICs

LAPIC
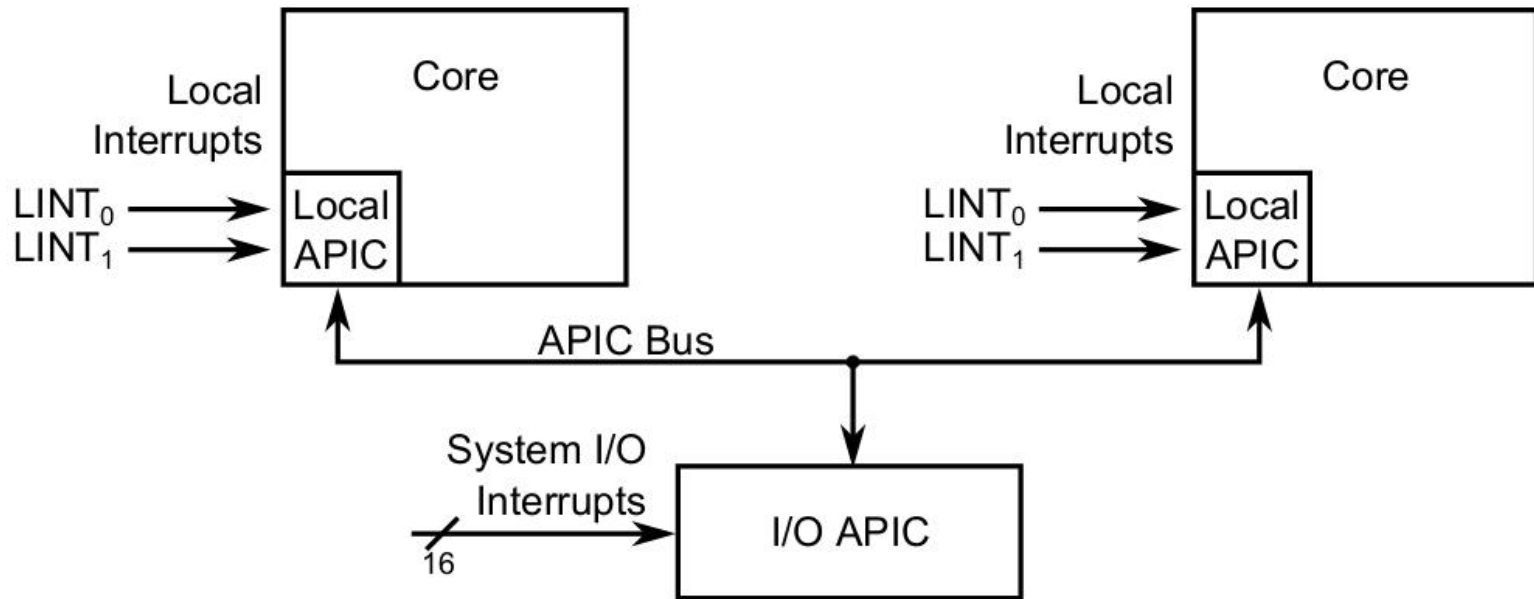
Interrupt Control Register

- The ICR register is used to initiate an IPI
- Values written into it specify the type of interrupt to be sent, and the target core

ICR (upper 32-bits)

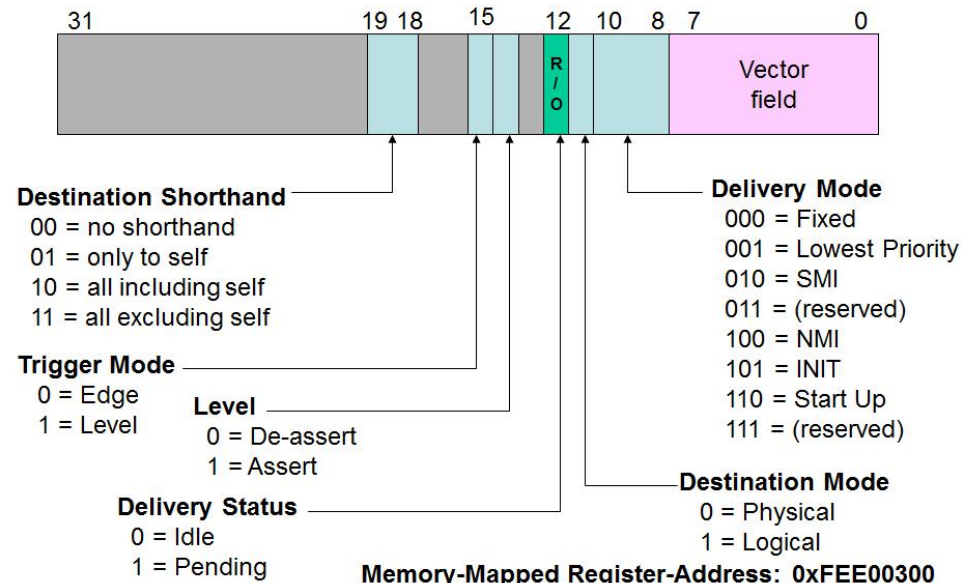| 31	24		0
Destination field		reserved	

The Destination Field (8-bits) can be used to specify which processor (or group of processors) will receive the message

Memory-Mapped Register-Address: 0xFEE00310

ICR (lower 32-bits)

31	19 18	15	12	10	8	7	0
			R/O			Vector field	

Destination Shorthand
00 = no shorthand
01 = only to self
10 = all including self
11 = all excluding self

Trigger Mode
0 = Edge
1 = Level

Level
0 = De-assert
1 = Assert

Delivery Status
0 = Idle
1 = Pending

Delivery Mode
000 = Fixed
001 = Lowest Priority
010 = SMI
011 = (reserved)
100 = NMI
101 = INIT
110 = Start Up
111 = (reserved)

Destination Mode
0 = Physical
1 = Logical

Memory-Mapped Register-Address: 0xFEE00300

Broadcast INIT-SIPI-SIPI Sequence

```
# address Local-APIC via register FS
    mov             $sel_fs, %ax
    mov             %ax, %fs


# broadcast 'INIT' IPI to 'all-except-self'
    mov      $0x000C4500, %eax ; 11 00 0 1 0 0 0 101 00000000
    mov      %eax, %fs:(0xFEE00300)
.B0: btl          $12, %fs:(0xFEE00300)
    jc    .B0


# broadcast 'Startup' IPI to 'all-except-self'
# using vector 0x11 to specify entry-point
# at real memory-address 0x00011000
    mov      $0x000C4611, %eax ; 11 00 0 0 1 0 0 0 110 00010001
    mov      %eax,  %fs:(0xFEE00300)
.B1: btl        $12,  %fs:(0xFEE00300)
    jc  .B1
```