

Dealing with Concurrency in the Kernel

Advanced Operating Systems and Virtualization

Alessandro Pellegrini

A.Y. 2018/2019



SAPIENZA
UNIVERSITÀ DI ROMA

Concurrent and Preemptive Kernels

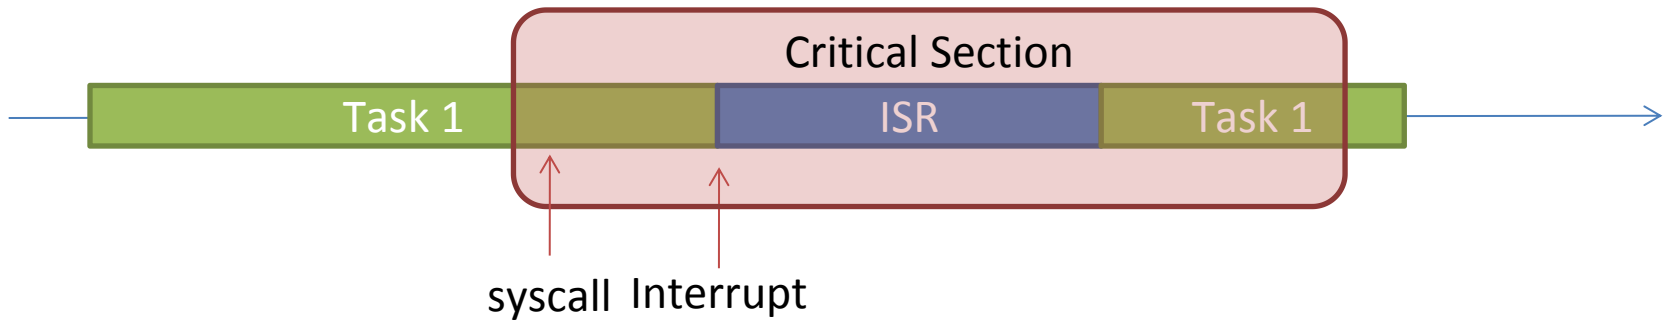
- Modern kernels are preemptive
 - A process running in kernel mode might be replaced by another process while in the middle of a kernel function
- Modern kernels run concurrently
 - Any core can run kernel functions at any time
- Kernel code must ensure consistency and avoid deadlock
- Typical solutions:
 - Explicit synchronization
 - Non-blocking synchronization
 - Data separation (e.g., per-CPU variables)
 - Interrupt disabling
 - Preemption disabling

} Mandatory on
multi-core machine



Kernel Race Conditions

- System calls and interrupts



```
queue_t *shared_q;
```

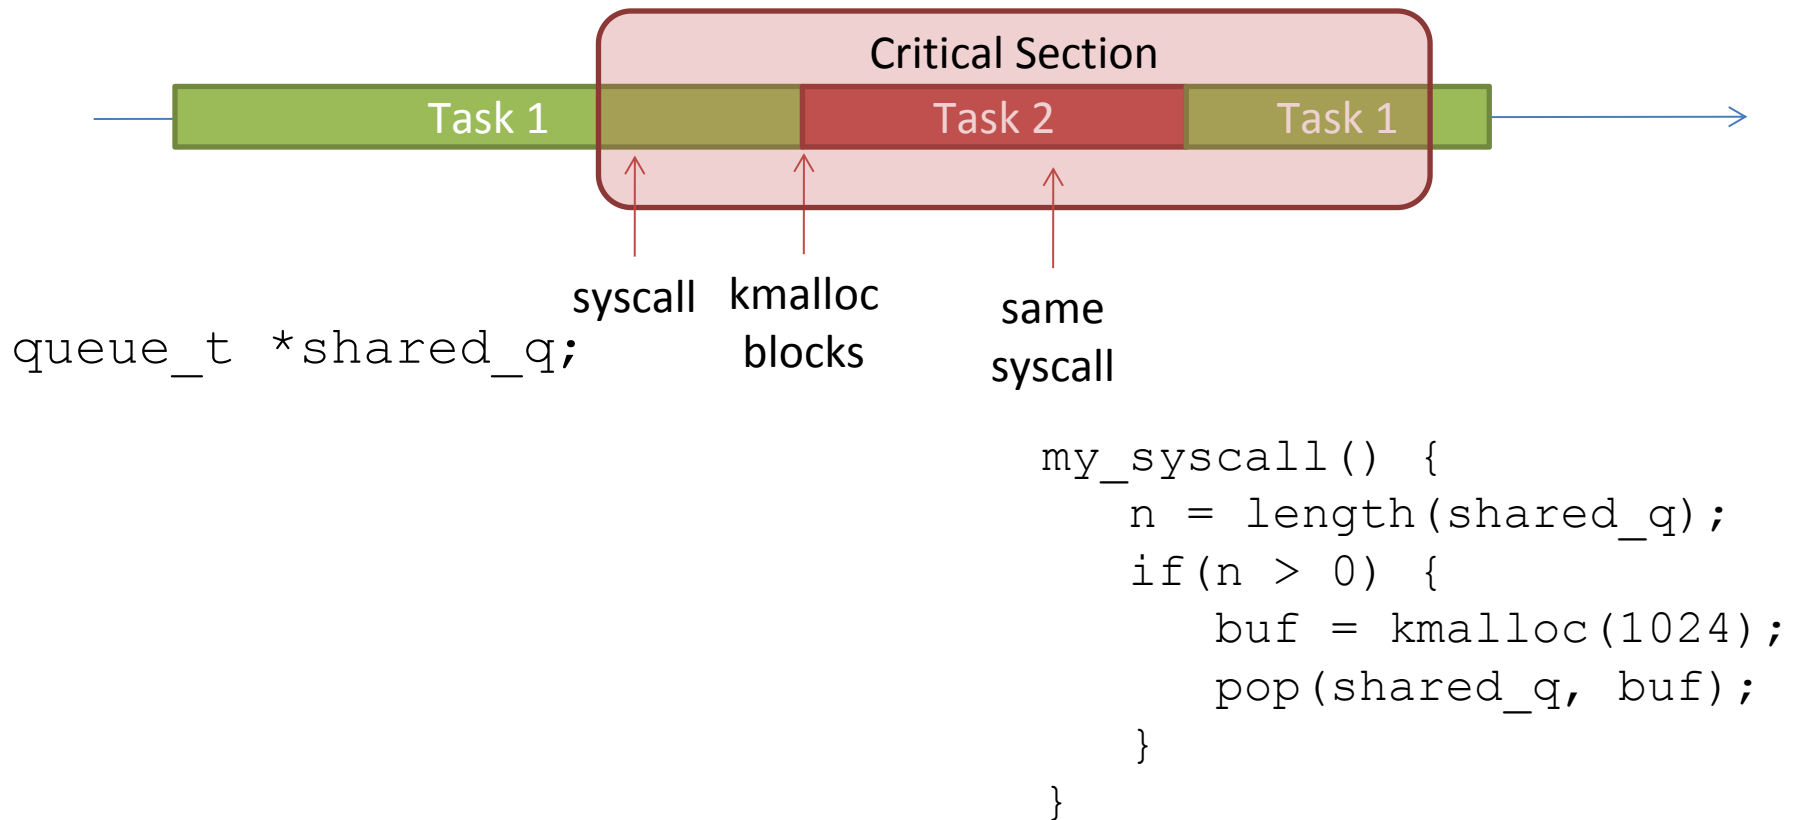
```
my_irq_handler() {  
    data = io(...);  
    push(shared_q, data);  
}
```

```
my_syscall() {  
    n = length(shared_q);  
    if(n > 0) {  
        buf = kmalloc(1024);  
        pop(shared_q, buf);  
    }  
}
```



Kernel Race Conditions

- System calls and preemption



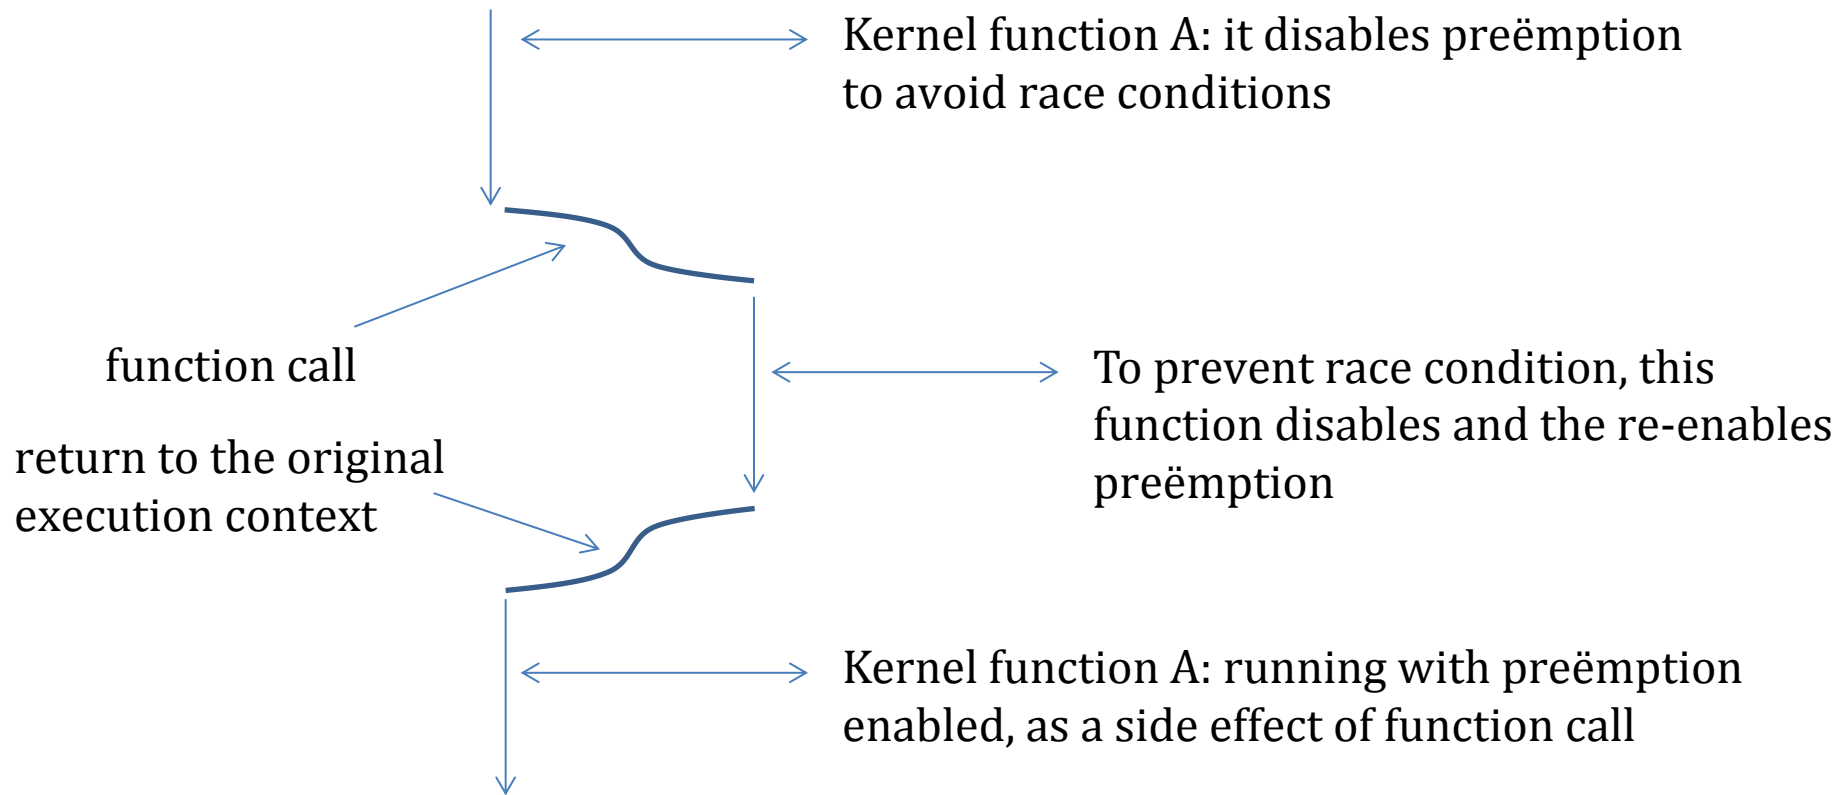
Enabling/Disabling Preemption

- Kernel preemption might take place when the scheduler is activated
- There must be a way to disable preemption
 - This is based on a (per-CPU) counter
 - A non-zero counter tells that preemption is disabled
- `preempt_count()`: return the current's core counter
- `preempt_disable()`: increases by one the preemption counter (needs a memory barrier).
- `preempt_enable()`: decreases by one the preemption counter (needs a memory barrier).



Why do we need counters?

- In a Kernel with no preemption counters this is possible:




Enabling/Disabling HardIRQs

- Given the per-CPU management of interrupts, HardIRQs can be disabled only locally
- Managing the IF flags:
 - `local_irq_disable()`
 - `local_irq_enable()`
 - `irqs_disabled()`
- Nested activations (same concept as in the preemption case):
 - `local_irq_save(flags)`
 - `local_irq_restore(flags)`



The _save Version

```
#define raw_local_irq_save(flags)          \  
do {                                       \  
    typecheck(unsigned long, flags);    \  
    flags = arch_local_irq_save();       \  
} while (0)
```



```
extern inline unsigned long native_save_fl(void)  
{  
    unsigned long flags;  
    asm volatile("pushf ; pop %0"  
        : "=rm" (flags)  
        : /* no input */  
        : "memory");  
    return flags;  
}
```

Why cannot we rely on
counters as in the case of
preemption disabling?



Per-CPU Variables

- A support to implement “data separation” in the kernel
- It is the best “synchronization” technique
 - It removes the need for explicit synchronization
- They are not silver bullets
 - No protection against asynchronous functions
 - No protection against preemption and reschedule on another core



Atomic Operations

- Based on RMW instructions
- `atomic_t` type
 - `atomic_fetch_{add,sub,and,andnot,or,xor}()`
- `DECLARE_BITMAP()` macro
 - `set_bit()`
 - `clear_bit()`
 - `test_and_set_bit()`
 - `test_and_clear_bit()`



Memory Barriers

- A compiler might reorder the instructions
 - Typically done to optimize the usage of registers
- Out of order pipeline and Memory Consistency models can reorder memory accesses
- Two families of barriers:
 - *Optimization* barriers
 - `#define barrier() asm volatile("":::"memory");`
 - *Memory* barriers
 - `{ smp_ } mb () : full memory barrier`
 - `{ smp_ } rmb () : read memory barrier`
 - `{ smp_ } wmb () : write memory barrier`

} Add fences
if necessary



Big Kernel Lock

- Traditionally called a "Giant Lock"
- This is a simple way to provide concurrency to userspace avoiding concurrency problems in the kernel
- Whenever a thread enters kernel mode, it acquires the BKL
 - No more than one thread can live in kernel space
- Completely removed in 2.6.39



Linux Mutexes

```
DECLARE_MUTEX(name);  
/* declares struct semaphore <name> ... */  
  
void sema_init(struct semaphore *sem, int val);  
/* alternative to DECLARE_... */  
  
void down(struct semaphore *sem); /* may sleep */  
  
int down_interruptible(struct semaphore *sem);  
/* may sleep; returns -EINTR on interrupt */  
  
int down_trylock(struct semaphore *sem);  
/* returns 0 if succeeded; will not sleep */  
  
void up(struct semaphore *sem);
```



Linux Spinlocks

```
#include <linux/spinlock.h>

spinlock_t my_lock = SPINLOCK_UNLOCKED;
spin_lock_init(spinlock_t *lock);
spin_lock(spinlock_t *lock);
spin_lock_irqsave(spinlock_t *lock, unsigned long flags);
spin_lock_irq(spinlock_t *lock);
spin_lock_bh(spinlock_t *lock);

spin_unlock(spinlock_t *lock);
spin_unlock_irqrestore(spinlock_t *lock,
                      unsigned long flags);
spin_unlock_irq(spinlock_t *lock);
spin_unlock_bh(spinlock_t *lock);
spin_is_locked(spinlock_t *lock);
spin_trylock(spinlock_t *lock)
spin_unlock_wait(spinlock_t *lock);
```



Linux Spinlocks

```
static inline void __raw_spin_lock_irq(raw_spinlock_t *lock)
{
    local_irq_disable();
    preempt_disable();
    spin_acquire(&lock->dep_map, 0, 0, _RET_IP_);
}
```

spin_lock_irq

spin_lock

preempt_disable

local_irq_disable

spin_lock_irqsave

spin_lock

preempt_disable

local_irq_save



Read/Write Locks

Read

Get Lock:

- Lock r
- Increment c
- if $c == 1$
 - lock w
- unlock r

Release Lock:

- Lock r
- Decrement c
- if $c == 0$
 - unlock w
- unlock r

Write

Get Lock:

- Lock w

Release Lock:

- Unlock w



Read/Write Locks

```
rwlock_t xxx_lock = __RW_LOCK_UNLOCKED(xxx_lock);  
unsigned long flags;
```

```
read_lock_irqsave(&xxx_lock, flags);  
.. critical section that only reads the info ...  
read_unlock_irqrestore(&xxx_lock, flags);
```

```
write_lock_irqsave(&xxx_lock, flags);  
.. read and write exclusive access to the info ...  
write_unlock_irqrestore(&xxx_lock, flags);
```

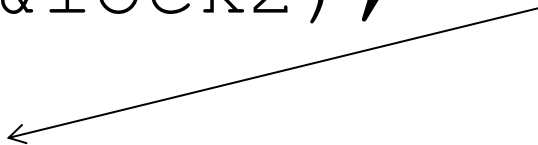



seqlocks

- A seqlock tries to tackle the following situation:
 - A small amount of data is to be protected.
 - That data is simple (no pointers), and is frequently accessed.
 - Access to the data does not create side effects.
 - It is important that writers not be starved for access.
- It is a way to avoid readers to starve writers



seqlocks

- `#include <linux/seqlock.h>`
- `seqlock_t lock1 =`
`SEQLOCK_UNLOCKED;`
- `seqlock_t lock2;`
- `seqlock_init(&lock2);`


Exclusive access and
increment the
sequence number
- `write_seqlock(&the_lock);`
- `/* Make changes here */`


increment again
- `write_sequnlock(&the_lock);`



seqlocks

- Readers do not acquire a lock:

```
unsigned int seq;
```

```
do {
```

```
    seq = read_seqbegin(&the_lock);
```

```
    /* Make a copy of the data of interest */
```

```
} while read_seqretry(&the_lock, seq);
```

- The call to `read_seqretry` checks whether the initial number was odd
- It additionally checks if the sequence number has changed



Read-Copy-Update (RCU)

- Another synchronization mechanism, added in October 2002
- RCU ensures that reads are coherent by maintaining multiple versions of objects and ensuring that they are not freed up until all pre-existing read-side critical sections complete
- RCU allow many readers and many writers to proceed concurrently
- RCU is lock-free (no locks nor counters are used)
 - Increased scalability, no cache contention on synchronization variables



Read-Copy-Update (RCU)

- Three fundamental mechanisms:
 - Publish-subscribe mechanism (for insertion)
 - Wait for pre-existing RCU readers to complete (for deletion)
 - Maintain multiple versions of RCU-updated objects (for readers)
- RCU scope:
 - Only dynamically allocated data structures can be protected by RCU
 - No kernel control path can sleep inside a critical section protected by RCU



Insertion

```
struct foo {  
    int a;  
    int b;  
    int c;  
};  
struct foo *gp = NULL;
```

```
/* . . . */
```

```
p = kmalloc(sizeof(*p), GFP_KERNEL);
```

```
p->a = 1;
```

```
p->b = 2;
```

```
p->c = 3;
```

```
gp = p;
```

Is this always correct?



Insertion

```
struct foo {  
    int a;  
    int b;  
    int c;  
};  
struct foo *gp = NULL;
```

```
/* . . . */
```

```
p = kmalloc(sizeof(*p), GFP_KERNEL);  
p->a = 1;  
p->b = 2;  
p->c = 3;
```

rcu_assign_pointer(gp, p) ← the "publish" part



Reading

```
p = gp;  
if (p != NULL) {  
    do_something_with(p->a, p->b, p->c);  
}
```

Is this always correct? ←



Reading

```
rcu_read_lock();
```

```
p = rcu_dereference(gp);
```

← Memory barriers here

```
if (p != NULL) {
```

```
    do_something_with(p->a, p->b, p->c);
```

```
}
```

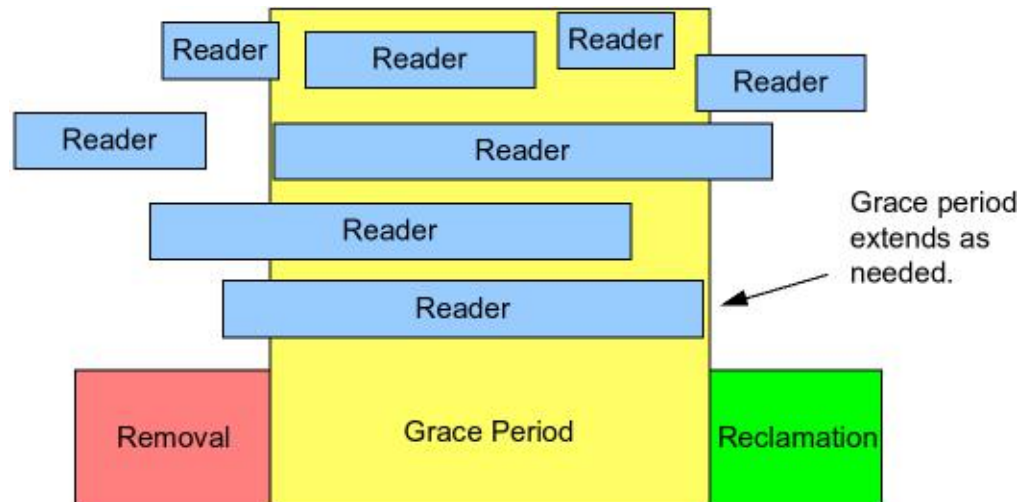
```
rcu_read_unlock();
```



Wait Pre-Existing RCU Updates

- `synchronize_rcu()`
- It can be schematized as:

```
for_each_online_cpu(cpu)  
    run_on(cpu);
```



Wait Pre-Existing RCU Updates

```
struct foo {
    struct list_head list;
    int a;
    int b;
    int c;
};

LIST_HEAD(head);

/* . . . */

p = search(head, key);
if (p == NULL) {
    /* Take appropriate action, unlock, and return. */
}

q = kmalloc(sizeof(*p), GFP_KERNEL);
*q = *p;
q->b = 2;
q->c = 3;
list_replace_rcu(&p->list, &q->list);
synchronize_rcu();
kfree(p);
```



Multiple Concurrent RCU Updates

```
struct foo {  
    struct list_head list;  
    int a;  
    int b;  
    int c;  
};  
LIST_HEAD(head);  
  
/* . . . */
```

```
p = search(head, key);  
if (p == NULL) {  
    /* Take appropriate action, unlock, and return. */  
}
```

```
q = kmalloc(sizeof(*p), GFP_KERNEL);  
*q = *p;  
q->b = 2;  
q->c = 3;  
list_replace_rcu(&p->list, &q->list);  
synchronize_rcu();  
kfree(p);
```

```
p = search(head, key);  
if (p != NULL) {  
    list_del_rcu(&p->list);  
    synchronize_rcu();  
    kfree(p);  
}
```



RCU Garbage Collection

- An old version of a data structure can be still accessed by readers
 - It can be freed only after that all readers have called `rcu_read_unlock()`
- A writer cannot waste too much time waiting for this condition
- `call_rcu()` registers a callback function to free the old data structure
- Callbacks are activated by a dedicated SoftIRQ action

